Selective Learning for Deep Time Series Forecasting

Yisong Fu[1,2], **Zezhi Shao**[1], **Chengqing Yu**[1,2], **Yujie Li**[1,2], **Zhulin An**[1,2],
Qi Wang[3], **Yongjun Xu**[1,2*], **Fei Wang**[1,2*]

[1]State Key Laboratory of AI Safety, Institute of Computing Technology,
Chinese Academy of Sciences   [2]University of Chinese Academy of Sciences
[3]Department of Automation, Tsinghua University

`{fuyisong24s, shaozezhi, yuchengqing22b, liyujie23s, anzhulin}@ict.ac.cn`,
`cheemswang@mail.tsinghua.edu.cn`, `{xyj,wangfei}@ict.ac.cn`

Abstract

Benefiting from high capacity for capturing complex temporal patterns, deep learning (DL) has significantly advanced time series forecasting (TSF). However, deep models tend to suffer from severe overfitting due to the inherent vulnerability of time series to noise and anomalies. The prevailing DL paradigm uniformly optimizes all timesteps through the MSE loss and learns those uncertain and anomalous timesteps without difference, ultimately resulting in overfitting. To address this, we propose a novel selective learning strategy for deep TSF. Specifically, selective learning screens a subset of the whole timesteps to calculate the MSE loss in optimization, guiding the model to focus on generalizable timesteps while disregarding non-generalizable ones. Our framework introduces a dual-mask mechanism to target timesteps: (1) an uncertainty mask leveraging residual entropy to filter uncertain timesteps, and (2) an anomaly mask employing residual lower bound estimation to exclude anomalous timesteps. Extensive experiments across eight real-world datasets demonstrate that selective learning can significantly improve the predictive performance for typical state-of-the-art deep models, including 37.4% MSE reduction for Informer, 8.4% for TimesNet, and 6.5% for iTransformer.

Code: https://github.com/GestaltCogTeam/selective-learning
https://github.com/GestaltCogTeam/BasicTS

1 Introduction

Time series forecasting (TSF) plays a crucial role in many real-world applications, such as traffic flow prediction [26, 45, 43], weather forecasting [38, 65, 80, 12], and energy consumption planning [33, 57]. The rapid advancement of deep learning (DL) has spurred breakthroughs in TSF, with numerous deep models pushing the boundaries of predictive performance and becoming pivotal in the field [81, 64, 82, 76, 35, 29].

Despite the strong capacity for capturing complex temporal patterns, deep TSF models are prone to suffer from severe overfitting issues under certain scenarios due to the characteristics of real-world time series data [42, 58, 7]. Unlike other data modalities such as natural language and images, time series is inherently susceptible to *noise* and *anomalies* introduced by random exogenous factors [28, 56, 31]. For example, industrial sensors are easily affected by noise from mechanical vibrations and electromagnetic disturbances, and stock prices exhibit non-stationary fluctuations given the policy interventions. These interference factors are challenging to model and typically change over time, exhibiting *uncertain* and *anomalous* patterns at a specific range of timesteps. However, the current
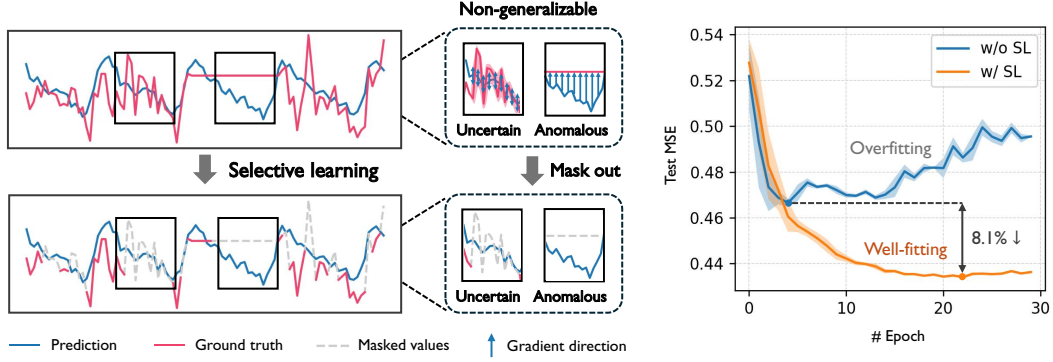
---

*Corresponding authors.

Figure 1: **Left:** When optimizing the model through MSE loss, our proposed selective learning calculates MSE only on a subset of timesteps, while masking out *uncertain* and *anomalous* ones that are *non-generalizable*. **Right:** Test MSE curves of iTransformer during training on the ETTh1 dataset (prediction length $F = 336$). The model exhibits severe overfitting, but this is effectively mitigated through selective learning, yielding an 8.1% reduction in test MSE with stable convergence.

DL paradigm treats each time step equally when training the models with regression loss functions (e.g., MSE/MAE loss). This causes an overfitting issue when forcing models to learn *uncertain* and *anomalous* timesteps unavailable to generalize, deteriorating models' performance. For example, iTransformer [29] encounters significant overfitting when trained on the ETTh1 dataset. As shown in Figure 1 (right), its test MSE during training gradually increases after the third epoch.

To address these challenges, we propose selective learning, a novel learning strategy for deep TSF. As illustrated in Figure 1 (left), the main idea of selective learning is to involve only generalizable timesteps, a subset of the time series, in optimization and discard identified *uncertain* or *anomalous* ones. In implementation, we propose a dual-mask mechanism to filter out *non-generalizable* timesteps dynamically. (1) For *uncertain* timesteps, we introduce the entropy of the prediction residual distribution as the uncertainty measure. With the sliding window sampling in time series, we can obtain multiple samples of a predicted timestep under different historical windows, thereby quantifying the entropy of the residual to serve as an indicator to filter out high-entropy ones. (2) For *anomalous* timesteps, we train an estimation model to obtain the residual lower bound of each timestep. By masking timesteps where current residuals are closest to the lower bound, it removes *non-generalizable* anomalies dynamically while keeping to-be-learned timesteps.

Extensive experiments across eight real-world datasets show that selective learning achieves consistent performance gains on six well-acknowledged deep models. It proves particularly effective for models that are susceptible to overfitting, where it achieves a **37.4%** reduction in MSE and for Informer [81], and **15.6%** in MSE for Crossformer [79]. Notably, selective learning maintains its benefits even for state-of-the-art baselines, such as TimesNet [63] (**8.4%** MSE reduction) and iTransformer [29] (**6.5%** MSE reduction). Furthermore, we conducted comparative analyses with alternative training objectives in §5.4. The consistent leading performance of selective learning further demonstrates the advantage of optimizing *generalizable* subsets over global sequence optimization.

In summary, our contribution is three-fold:

- We propose selective learning, a novel learning strategy for deep TSF to identify and expel *non-generalizable* timesteps in optimization. It is the first trial to address the overfitting issue from the timestep granularity in the field.

- Technically, we devise a tractable strategy by introducing a dual-mask mechanism to filter out *uncertain* and *anomalous* timesteps dynamically during training.

- Our method is agnostic to deep learning backbones and examined across several real-world datasets. The results demonstrate its effectiveness and show consistent performance improvement over all baselines.

2 Related Work

2.1 Deep Models for Time Series Forecasting

In recent years, numerous deep models have been proposed to capture complex dependencies in TSF. Transformer-based models have gained significant attention for their ability to capture long-term temporal dependencies through attention mechanisms [81, 64, 82, 74, 79, 35, 29]. For example, Informer [81] introduces a ProbSparse attention to reduce the quadratic complexity. PatchTST [35] splits time series into patches and employs a channel-independent strategy. iTransformer [29] embeds each series independently to the variate token and applies self-attention to capture multivariate correlations. In contrast to Transformers, CNN-based models [63, 32, 25] exhibit strong proficiency in extracting local patterns. Typically, TimesNet [63] transforms time series into 2D tensors and employs CNN to capture inter- and intra-period dependencies. Additionally, MLP-based models offer efficient alternatives with lightweight architectures [44, 13, 55, 70, 76]. DLinear [76] leverages a simple linear layer with decomposition, and TimeMixer [55] captures multi-scale information through MLP layers. Our proposed selective learning can be easily applied to these deep models.

2.2 Training Strategies for Time Series Forecasting

The prevailing DL paradigm computes the regression loss (e.g., MSE/MAE) uniformly across all timesteps, and some works have explored alternative training strategies. For example, iTransformer [29] proposes a training strategy that randomly selects subsets of variables for large-scale multivariate time series. Merlin [75] employs a knowledge distillation [15] strategy to enhance the model's robustness against data missing. These approaches are optimized for specific scenarios or tightly coupled with model architectures, limiting their broader applicability. The most relevant work is MTGNN [67], which applies the idea of curriculum learning [4] to TSF by progressively increasing the prediction length during training. However, it overlooks that difficulty is not solely determined by prediction length but is also influenced by intrinsic data characteristics. Selective learning addresses this issue by masking *non-generalizable* timesteps while maintaining broad applicability.

Another line of work proposes alternative training objectives to replace the regression loss. For example, Soft-DTW [8], DILATE [23], and TILDE-Q [24] align the shape between predictions and target sequences under temporal distortions. FreDF [51] combines the MSE loss with a frequency loss, mitigating the label correlation. PS loss [21] enhances the alignment by incorporating patch-wise distribution information. These works focus on matching the shape or distribution in temporal or frequency domain between sequences, but none have recognized that global alignment over the whole sequence is not optimal, as certain timesteps in the target sequence are inherently *non-generalizable*.

3 Preliminaries

Notations For a multivariate time series with N variables, let $X_t \in \mathbb{R}^N$ represent the t-th timestep. Given a historical time series $\mathbf{X}_{t-L:t} = \{X_{t-L}, X_{t-L+1}, \cdots, X_{t-1}\} \in \mathbb{R}^{L \times N}$, where L is the look-back window size, the TSF task is to predict future values $\hat{\mathbf{X}}_{t:t+F} = \{X_t, X_{t+1}, \cdots, X_{t+F-1}\} \in \mathbb{R}^{F \times N}$ with forecasting window size F. Considering a historical time series $\mathbf{X}_{0:T} \in \mathbb{R}^{T \times N}$ for training, the training dataset $\mathcal{D}_{train} = \{(\mathbf{X}_{t-L:t}, \mathbf{X}_{t:t+F})\}_{t=L}^{T-F}$ is constructed by a sliding window approach with stride 1.

Problem Statement The current DL paradigm is to find the best mapping from the samples in $\mathcal{D}_{train}$, i.e., $\hat{\mathbf{X}}_{t:t+F} = f(\mathbf{X}_{t-L:t}; \boldsymbol{\theta})$, where $f(\cdot; \boldsymbol{\theta}) : \mathbb{R}^{L \times N} \to \mathbb{R}^{F \times N}$ is a deep neural network parameterized by θ. Mean squared error (MSE) measures the discrepancy between the prediction $\hat{\mathbf{X}}_{t:t+F}$ and the ground truth $\mathbf{X}_{t:t+F}$ and is one of the commonly used loss functions to optimize $\boldsymbol{\theta}$:

$$\mathcal{L}_{MSE}(\boldsymbol{\theta}) = \frac{1}{N \cdot F} \sum_{i=0}^{F-1} ||X_{t+i} - f(\mathbf{X}_{t-L:t}; \boldsymbol{\theta})_i||^2, \tag{1}$$

$$\boldsymbol{\theta}_{\tau+1} = \boldsymbol{\theta}_\tau - \eta \nabla_{\boldsymbol{\theta}} \mathcal{L}_{MSE}, \tag{2}$$

where η is the learning rate. We use τ to denote the number of iterations during training, distinguishing it from the timestep index t.
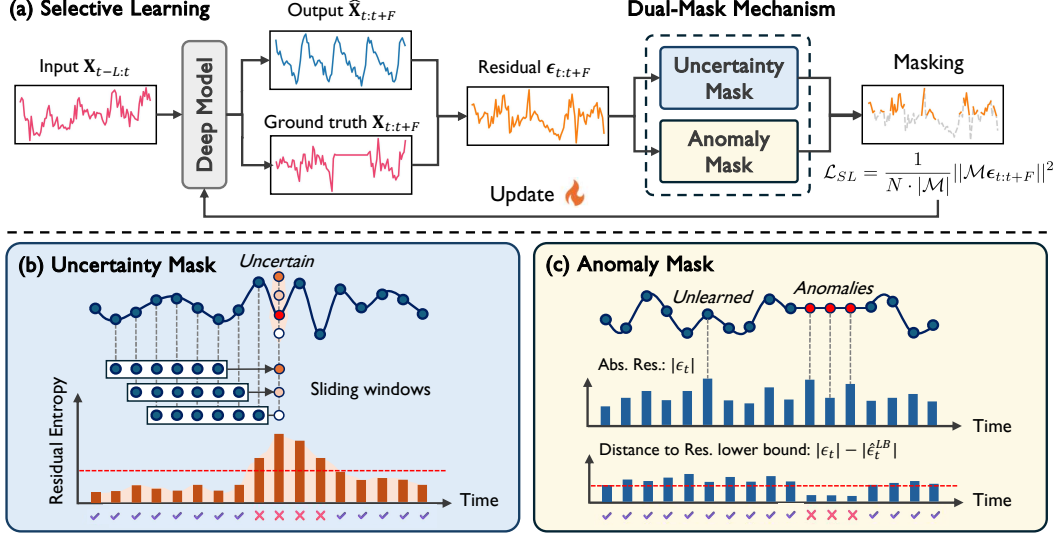
3

Figure 2: (a) Overall framework of selective learning. (b) Uncertainty mask. (c) Anomaly mask.

4 Selective Learning

4.1 Overview of Selective Learning

We propose selective learning, a model-agnostic learning strategy for deep TSF to address the overfitting issue. The main idea of selective learning is to calculate the MSE loss only on a subset of timesteps. This enables models to focus selectively on *generalizable* timesteps while disregarding *non-generalizable* ones. Figure 2 (a) illustrates the overall framework of selective learning. The implementation details and workflow of selective learning are provided in Appendix C.3.

As illustrated in Figure 1 (left), we identify two critical categories of timesteps that degrade model generalizability. (1) *Uncertain* timesteps. Primarily originating from inherent noise in time series (e.g., signal disturbances [23]), they are characterized by high predictive uncertainty. Consequently, the gradients at these timesteps update towards random directions, resulting in undesirable fitting to the noise. (2) *Anomalous* timesteps. They are mainly caused by exogenous exceptional events (e.g., sensor malfunctions [73, 72]). The model's predictions, though possibly confident, exhibit significant errors. This forces the model to learn instance-specific features through biased gradient updates, ultimately harming generalization.

To this end, we propose a dual-mask mechanism to dynamically filter out these non-generalizable timesteps, as shown in Figure 2. Formally, we define selective learning as follows: Given a deep TSF model $f(\cdot, \boldsymbol{\theta})$ at τ-th iteration, we find a mask $\mathcal{M}^{(\tau)} \in \{0,1\}^F$ that constrains optimization only over a subset of timesteps:

$$\mathcal{L}_{SL}(\boldsymbol{\theta}) = \frac{1}{N \cdot |\mathcal{M}^{(\tau)}|} \sum_{i=0}^{F-1} ||\mathcal{M}^{(\tau)}(X_{t+i} - f(\mathbf{X}_{t-L:t}; \boldsymbol{\theta})_i)||^2, \qquad (3)$$

$$\boldsymbol{\theta}_{\tau+1} = \boldsymbol{\theta}_\tau - \eta \nabla_{\boldsymbol{\theta}} \mathcal{L}_{SL}, \qquad (4)$$

where $\mathcal{M}^{(\tau)} = \mathcal{M}_u^{(\tau)} \vee \mathcal{M}_a^{(\tau)}$, and $\vee$ is the element-wise OR operator. $\mathcal{M}_u^{(\tau)}, \mathcal{M}_a^{(\tau)} \in \{0,1\}^F$ are the uncertainty mask and anomaly mask, respectively. We will describe them in detail in the following sections. Since uncertain and anomalous patterns do not necessarily appear synchronously across all variables, we adopt the channel-independent strategy [35], generating masks for each variable independently. For theoretical tractability, we will focus on the univariate case in subsequent analysis, with natural extensibility to multivariate scenarios due to channel independence.

4.2 Uncertainty Mask

We propose an entropy-based uncertainty masking approach for those timesteps that exhibit high predictive uncertainty. Let $\epsilon_t = X_t - \hat{X}_t$ denote the residual of t-th timesteps. The differential entropy of ϵ_t is

$$H(\epsilon_t) = \int p(\epsilon_t) \ln p(\epsilon_t) \mathrm{d}t. \tag{5}$$

Since training samples are constructed via sliding windows over $\mathbf{X}_{1:T}$, each timestep will be predicted n_t times in one epoch, where $n_t = \min\{t - L + 1, F\}$. Therefore, we can estimate the residual distribution of ϵ_t using the most recent n_t predictions. Let $l(\boldsymbol{\psi}|\epsilon_t)$ be the likelihood model of the residual, and we have

$$\hat{\boldsymbol{\psi}} = \arg\max_{\boldsymbol{\psi}} l(\boldsymbol{\psi}|\epsilon_t^{(1)}, \epsilon_t^{(2)}, \cdots, \epsilon_t^{(n_t)}), \tag{6}$$

$$\hat{H}(\epsilon_t) = \int l(\hat{\boldsymbol{\psi}}|\epsilon_t) \ln l(\hat{\boldsymbol{\psi}}|\epsilon_t) \mathrm{d}t, \tag{7}$$

where $\epsilon_t^{(i)}$ is the i-th most recent prediction residual.

In practice, we assume that the residual $\epsilon_t \sim \mathcal{N}(\mu_t, \sigma_t^2)$, therefore we have

$$\hat{H}(\epsilon_t) = \frac{1}{2} \ln(2\pi e \hat{\sigma}_t^2), \tag{8}$$

$$\hat{\sigma}_t^2 = \frac{1}{n_t} \sum_{i=1}^{n_t} (\epsilon_t^{(i)} - \bar{\epsilon}_t)^2 \tag{9}$$

Since these residuals $\epsilon_t^{(i)}$ are computed at different training time τ, they originate from distinct $f(\cdot, \boldsymbol{\theta}_\tau)$. Under Assumptions 2-4 (Appendix A.1), Theorem 1 provides an upper bound for the error introduced by different $\boldsymbol{\theta}$.

Theorem 1 (Upper Bound for Variance Estimation Error). *The error bound between variance estimation under distinct parameters $\hat{\sigma}_t^2$ and that under identical parameters $\hat{\sigma}_t^2(\boldsymbol{\theta}_\tau)$ satisfies:*

$$|\hat{\sigma}_t^2 - \hat{\sigma}_t^2(\boldsymbol{\theta}_\tau)| \leq 4L_f R\eta G(2K - 1), \tag{10}$$

where K is the number of iterations per epoch, and L_f, R, G are constants.

The proof is provided in Appendix A. According to Theorem 1, we can always control the estimation error by choosing a sufficiently small learning rate η and a large batch size.

We employ a hard thresholding γ_u on the top-$r_u\%$ residual entropy and obtain the uncertainty mask $\mathcal{M}_u^{(\tau)} \in \{0, 1\}^F$ satisfying

$$(\mathcal{M}_u^{(\tau)})_t = \begin{cases} 0, & \hat{H}(\epsilon_t) > \gamma_u, \\ 1, & \text{otherwise.} \end{cases} \tag{11}$$

4.3 Anomaly Mask

Predictions for anomalous timesteps typically exhibit significantly larger residuals due to deviations in ground truth values. The most intuitive solution is to mask timesteps with high $|\epsilon_t|$ to filter anomalies. However, this naive approach suffers from a critical limitation: it indiscriminately excludes both genuine anomalies and currently unlearned yet (but potentially generalizable) patterns, particularly during the early stage of training when the learning process remains incomplete.

To overcome this limitation, we draw inspiration from practices in other fields [34, 48, 27, 50], and define $S(X_t)$ as the deviation between the residual and its theoretical lower bound ϵ_t^{LB}:

$$S(X_t) = |X_t - f(\mathbf{X}; \boldsymbol{\theta})_t)| - \epsilon_t^{LB}. \tag{12}$$

This formulation enables a key separation: Anomalous timesteps exhibit elevated residual lower bounds, resulting in comparatively small $S(X)$ values. In contrast, unlearned timesteps demonstrate larger $S(X)$ values due to significant gaps between current residuals and their theoretical minima.

In practice, we train a lightweight model $g(\cdot; \phi)$ on $\mathcal{D}_{train}$ to estimate the residual lower bound, thereby estimating $S(X_t)$ by:

$$\hat{S}(X_t) = \underbrace{|X_t - f(\mathbf{X}; \boldsymbol{\theta})_t|}_{\text{residual } \epsilon_t} - \underbrace{|X_t - g(\mathbf{X}; \boldsymbol{\phi})_t|}_{\text{estimated LB } \hat{\epsilon}_t^{LB}}. \quad (13)$$

The details of the estimation model are discussed in Appendix C.3. Analogous to the uncertainty mask, we employ a hard thresholding γ_a to filter out the top-r_u% of timesteps with the smallest $\hat{S}(X_t)$ and obtain the anomaly mask $\mathcal{M}_a^{(\tau)} \in \{0, 1\}^F$:

$$(\mathcal{M}_a^{(\tau)})_t = \begin{cases} 0, & \hat{S}(X_t) < \gamma_a, \\ 1, & \text{otherwise.} \end{cases} \quad (14)$$

Notably, instead of using the estimated residual lower bound as a static masking criterion, $S(X_t)$ can dynamically adjust the masking based on the current predictions. This approach offers two key advantages: (1) Static masking significantly alters the distribution of $\mathcal{D}_{train}$, thereby introducing bias, whereas dynamic masking adapts the mask during training to mitigate this in expectation. (2) For rare but critical extreme events (e.g., extreme weather) [10, 78] that are less generalizable, dynamic masking first learns the most generalizable timesteps and gradually attempts to learn timesteps previously considered anomalies. See Appendix E.2 for detailed discussions.

5   Experiments

5.1   Experimental Setup

Datasets   We thoroughly evaluate the effectiveness of the proposed selective learning on 8 real-world datasets, including Electricity, Exchange, Weather, ILI, and 4 ETT datasets (ETTh1, ETTh2, ETTm1, ETTm2), which have been extensively used for benchmarking [42, 64, 37]. A detailed description of the datasets is provided in Appendix C.1.

Baselines   Selective learning is a model-agnostic training strategy, and it is compatible with any deep TSF models. We carefully select six well-acknowledged deep models as the baselines, including Transformer-based models (Informer [81], Crossformer [79], PatchTST [35], iTransformer [29]), CNN-based models (TimesNet [63]), and MLP-based models (TimeMixer [55]). See Appendix C.2 for the introduction to the baselines. We select DLinear [76] as the estimation model for all baselines, and we further discuss the effects of different estimation models in §5.5.

Experimental Settings   All baselines follow the same experimental setup with prediction lengths $F \in \{24, 36, 48, 60\}$ for ILI and $F \in \{96, 192, 336, 720\}$ for others [63]. We search for the look-back window L and report the best results. For fair evaluation, when training baselines with selective learning to enhance their performance, we follow their original hyperparameter settings and only tune the masking ratios r_a and r_n. We utilize Adam [20] for the model optimization. We evaluate the performance of all baselines using two commonly used metrics, MSE and MAE. All experiments are implemented with PyTorch and conducted on 8 NVIDIA GeForce RTX 4090 24GB GPUs.

5.2   Main Results

Table 1 shows the forecasting results with and without selective learning. The results are averaged over three runs. The lower MSE/MAE indicates a more accurate prediction. Our comprehensive evaluations demonstrate that selective learning consistently enhances model performance in all 192 cases (see full results in Appendix G.1). Selective learning proves particularly impactful for early-generation architectures that are susceptible to overfitting, where it achieves an average reduction of **37.4%** in MSE and **25.4%** in MAE for Informer [81] (**66.8%** MSE and **42.6%** MAE reduction in the ETTm2 dataset), and **15.6%** in MSE and **10.5%** in MAE for Crossformer [79]. Notably, it maintains its benefits even for state-of-the-art baselines, such as iTransformer [29] (**6.5%** MSE and **4.4%** MAE reduction) and TimeMixer [55] (**4.3%** MSE and **3.3%** MAE reduction), where the improvements persist in models already equipped with RevIN [19]. This confirms that selective learning provides additional performance gains over existing distribution shift mitigation techniques.

6

Table 1: Comparison of forecasting results without/with selective learning (SL). We use prediction lengths $F \in \{24, 36, 48, 60\}$ for ILI and $F \in \{96, 192, 336, 720\}$ for other datasets. Results are averaged from all prediction lengths. Better results are in **bold**, and Δ denotes the improvements caused by selective learning. Full results of ETTh2 and ETTm1 are provided in Appendix G.1.

Method	ETTh1 MSE	ETTh1 MAE	ETTm2 MSE	ETTm2 MAE	Electricity MSE	Electricity MAE	Exchange MSE	Exchange MAE	Weather MSE	Weather MAE	ILI MSE	ILI MAE
Informer	1.289	0.917	1.485	0.919	0.342	0.420	1.520	0.985	0.337	0.374	4.953	1.544
+SL	**0.538**	**0.534**	**0.494**	**0.527**	**0.292**	**0.386**	**0.814**	**0.712**	**0.273**	**0.299**	**3.997**	**1.360**
Δ	-58.3%	-41.8%	-66.8%	-42.6%	-14.4%	-8.21%	-46.4%	-27.8%	-19.1%	-20.1%	-19.3%	-11.9%
Crossformer	0.455	0.465	0.588	0.528	0.182	0.277	0.755	0.649	0.226	0.284	3.982	1.342
+SL	**0.431**	**0.441**	**0.370**	**0.408**	**0.168**	**0.264**	**0.527**	**0.525**	**0.213**	**0.265**	**3.681**	**1.285**
Δ	-5.33%	-5.22%	-37.1%	-22.7%	-7.71%	-4.87%	-30.1%	-19.2%	-5.97%	-6.61%	-7.56%	-4.26%
PatchTST	0.427	0.433	0.271	0.329	0.167	0.262	0.342	0.396	0.228	0.262	2.076	0.921
+SL	**0.410**	**0.417**	**0.252**	**0.312**	**0.165**	**0.258**	**0.337**	**0.384**	**0.225**	**0.250**	**1.905**	**0.895**
Δ	-4.10%	-3.70%	-6.83%	-5.32%	-1.05%	-1.53%	-1.46%	-3.10%	-1.32%	-4.77%	-8.26%	-2.74%
TimesNet	0.499	0.486	0.289	0.343	0.198	0.301	0.382	0.429	0.248	0.284	2.493	1.028
+SL	**0.429**	**0.439**	**0.258**	**0.316**	**0.191**	**0.294**	**0.363**	**0.413**	**0.239**	**0.271**	**2.154**	**0.931**
Δ	-14.0%	-9.67%	-10.7%	-7.74%	-3.29%	-2.33%	-4.97%	-3.73%	-3.54%	-4.58%	-13.6%	-9.36%
iTransformer	0.458	0.457	0.273	0.332	0.164	0.257	0.364	0.413	0.235	0.269	1.909	0.914
+SL	**0.415**	**0.425**	**0.256**	**0.313**	**0.157**	**0.249**	**0.343**	**0.399**	**0.229**	**0.257**	**1.710**	**0.857**
Δ	-9.29%	-6.90%	-6.40%	-5.51%	-4.27%	-2.92%	-5.78%	-3.45%	-2.87%	-4.28%	-10.4%	-6.24%
TimeMixer	0.443	0.445	0.265	0.323	0.163	0.259	0.348	0.400	0.230	0.276	2.163	0.932
+SL	**0.411**	**0.421**	**0.251**	**0.309**	**0.160**	**0.254**	**0.335**	**0.394**	**0.226**	**0.268**	**2.026**	**0.895**
Δ	-7.11%	-5.40%	-5.01%	-4.26%	-2.15%	-2.12%	-3.60%	-1.38%	-1.74%	-2.81%	-6.37%	-4.06%

5.3 Zero-shot Forecasting

We conducted zero-shot forecasting experiments to evaluate the generalization benefits of selective learning across different datasets. Following prior works [83, 17, 5], we trained the models on dataset $\mathcal{D}_A$ and assessed on unseen dataset $\mathcal{D}_B$ without further training. As shown in Table 2, selective learning consistently enhance the performance of the baselines across diverse datasets in zero-shot forecasting, demonstrating its generalization advantage. Notably, in challenging generalization scenarios (ETTh2→ETTh1 and ETTm2→ETTm1), selective learning achieves significant improvements, with MSE reduced by **22.6%** and MAE by **14.5%** on average. Furthermore, in cases like ETTh1→ETTh2 and ETTm1→ETTm2, the results outperforms training from scratch on the target dataset, underscoring the benefits of selective learning.

Table 2: Zero-shot forecasting results on ETT datasets without/with selective learning. $\mathcal{D}_A \rightarrow \mathcal{D}_B$ denotes that the model was trained on $\mathcal{D}_A$ and tested on $\mathcal{D}_A$. The results are averaged from all prediction lengths. Better results are in **bold**, and red indicates a better result than training from scratch on $\mathcal{D}_B$ without selective learning.

Method	TimesNet w/o MSE	TimesNet w/o MAE	TimesNet +SL MSE	TimesNet +SL MAE	iTransformer w/o MSE	iTransformer w/o MAE	iTransformer +SL MSE	iTransformer +SL MAE	TimeMixer w/o MSE	TimeMixer w/o MAE	TimeMixer +SL MSE	TimeMixer +SL MAE
ETTh1→ETTh2	0.469	0.465	**0.419**	**0.439**	0.421	0.434	**0.394**	**0.420**	0.424	0.438	**0.389**	**0.416**
ETTh1→ETTm2	0.359	0.397	**0.351**	**0.389**	0.322	0.371	**0.310**	**0.359**	0.313	0.360	**0.300**	**0.352**
ETTh2→ETTh1	0.828	0.642	**0.595**	**0.517**	0.622	0.555	**0.501**	**0.487**	0.679	0.568	**0.560**	**0.510**
ETTm1→ETTh2	0.483	0.485	**0.465**	**0.469**	0.447	0.456	**0.428**	**0.444**	0.441	0.453	**0.434**	**0.445**
ETTm1→ETTm2	0.321	0.363	**0.288**	**0.339**	0.276	0.334	**0.271**	**0.326**	0.275	0.326	**0.267**	**0.319**
ETTm2→ETTm1	0.727	0.578	**0.459**	**0.448**	0.554	0.498	**0.443**	**0.439**	0.451	0.466	**0.428**	**0.425**

5.4 Comparison with Other Training Objectives

Our proposed selective learning exhibits strong compatibility. It is completely model-agnostic and can be applied to any deep learning architecture with various normalization methods [19, 11, 14]. Moreover, it maintains compatibility with learning strategies such as curriculum learning for TSF

Table 3: Comparison between selective learning (SL) and other training objectives with iTransformer as backbone. The results are averaged from all prediction lengths. The best results are in **bold**, and the second-best are underlined. Full results are provided in Appendix G.2.

Training objective	SL		PS		FreDF		TILDE-Q		MSE	
Metric	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE
ETTh1	**0.415**	**0.425**	<u>0.427</u>	0.440	0.450	0.455	0.432	<u>0.439</u>	0.458	0.457
ETTm2	**0.257**	**0.315**	0.264	0.320	<u>0.262</u>	<u>0.319</u>	0.263	<u>0.319</u>	0.273	0.332
Exchange	**0.343**	**0.399**	0.366	<u>0.409</u>	0.376	0.413	0.369	0.414	<u>0.364</u>	0.413
Weather	**0.229**	**0.257**	0.233	0.265	0.239	0.274	<u>0.232</u>	<u>0.262</u>	0.235	0.269

[67, 42]. However, selective learning operates on point-wise training objectives. In this section, we compare our selective learning with alternative non-point-wise training objectives, including shape-based (TILDE-Q [24]), frequency-based (FreDF [51]), and distribution-based (PS loss [21]) objectives. As shown in Table 3, selective learning consistently achieves superior performance. These results demonstrate that training the model by global alignment over whole sequences, whether in shape or distribution, in the temporal or frequency domain, proves suboptimal, validating the effectiveness of selective learning.

5.5 Ablation Study and Hyperparameter Analysis

Ablation Study To study the effectiveness of the components of selective learning, we conduct an ablation study covering: (1) removing either mask from the dual-mask mechanism, and (2) replacing the dual-mask mechanism with random masking with the same masking ratio. The results in Table 4 demonstrate that the model with full selective learning consistently achieves the best performance. Removing either mask leads to significant performance degradation across all four datasets, indicating that both masks contribute essential and distinct functionalities in filtering out non-generalizable patterns. Additionally, replacing the dual-mask mechanism with random masking reduces model performance to levels comparable to or worse than the unmasked counterparts. This suggests that randomly attending to a subset of timesteps fails to enhance the model's performance and generalizability. The effectiveness of selective learning fundamentally stems from our dual-mask mechanism, which directs model attention to *generalizable* timesteps while filtering out *non-generalizable* ones.

Table 4: Ablation results for selective learning with iTransformer as backbone. The results are averaged from all predicted lengths. Full ablation results are provided in Appendix G.3.

Dataset		ETTh1		ETTm2		Electricity		Weather	
Metric		MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE
Selective Learning		**0.415**	**0.425**	**0.257**	**0.315**	**0.157**	**0.249**	**0.229**	**0.257**
w/o	Uncertainty mask	0.436	0.443	0.265	0.322	0.162	0.256	0.232	0.266
	Anomaly mask	0.431	0.438	0.266	0.323	0.159	0.252	0.234	0.267
Replace	Random mask	0.457	0.460	0.274	0.332	0.165	0.261	0.237	0.269

Effects of Masking Ratio After validating the effectiveness of the dual-mask mechanism through ablation studies, we further investigate the effects of the masking ratios. When investigating a particular mask, we vary its masking ratio while fixing the other masking ratio at 0. The results are shown in Figure 3. We can observe that larger masking ratios demonstrate superior performance on highly non-stationary datasets (ETTh1 and Exchange). This indicates severe overfitting in such datasets, where models benefit from focusing selectively on the most generalizable patterns. In contrast, datasets exhibiting periodic patterns (Weather) show improved performance with smaller masking ratios. Besides, it can be observed that on the Exchange dataset, the 90% anomaly masking ratio yields peak performance. This occurs because market-induced non-generalizable anomalies in this dataset exert significantly greater influence than noise. However, in most scenarios, the best

Figure 3: Forecasting results under different masking ratios. The prediction length is 336.

results are achieved by a combination of two masking strategies, the ratios of which constitute critical hyperparameters of selective learning. Detailed guidelines for selecting optimal ratios are provided in the appendix C.4.

Effects of Estimation Model In the experiments above, we employed DLinear as the lightweight estimation model. In this section, we investigate the effects of different estimation models. We fix the iTransformer as the backbone model and additionally compare three estimation models: MLP, TimeMixer, and iTransformer. The results are shown in Figure 4. We can observe that: For highly non-stationary datasets (ETTh1 and Exchange), simpler models demonstrate superior performance; Conversely, datasets exhibiting periodic patterns (Weather) benefit from more complex estimation models. Despite this, the choice of the estimation model has a limited overall impact on performance, underscoring the robustness of the selective learning.


Figure 4: Forecasting performance with iTransformer as backbone and various estimation models. The results are averaged from all prediction lengths.

5.6 Learning Curve Analysis

In Figure 1, we initially illustrated iTransformer's training dynamics, highlighting selective learning's capacity to mitigate overfitting. To further validate this, Figure 5 presents additional learning curves on the ETTh1 dataset. While all three models exhibit varying degrees of overfitting, their counterparts trained with selective learning achieve stable convergence and superior performance. This demonstrates the efficacy of selective learning in mitigating overfitting and enhancing generalizability.


Figure 5: Test MSE curve on the ETTh1 dataset. The prediction length is 336.

9

6 Conclusion

In this work, we introduced selective learning, a novel strategy to mitigate overfitting in deep TSF by selectively computing regression loss on generalizable timesteps. Our dual-mask mechanism, comprising an uncertainty mask (based on residual entropy) and an anomaly mask (leveraging residual lower-bound estimation), dynamically filters non-generalizable timesteps, allowing models to focus on robust patterns. Extensive experiments across eight real-world datasets validate that selective learning improves predictive accuracy and model generalizability. The scope of this work is currently constrained to in-domain time series forecasting. Future work can investigate the generalization to diverse time series analysis tasks (e.g., classification, imputation) and explore pretraining strategies for time series foundation models. See Appendix F for limitations discussion and future directions.

Acknowledgement

References

[1] Abdul Fatir Ansari, Lorenzo Stella, Ali Caner Turkmen, Xiyuan Zhang, Pedro Mercado, Huibin Shen, Oleksandr Shchur, Syama Sundar Rangapuram, Sebastian Pineda Arango, Shubham Kapoor, et al. Chronos: Learning the language of time series. *Transactions on Machine Learning Research*, 2024.

[2] Sercan O Arik, Nathanael C Yoder, and Tomas Pfister. Self-adaptive forecasting for improved deep learning on non-stationary time-series. *arXiv preprint arXiv:2202.02403*, 2022.

[3] Tianyi Bai, Ling Yang, Zhen Hao Wong, Fupeng Sun, Xinlin Zhuang, Jiahui Peng, Chi Zhang, Lijun Wu, Qiu Jiantao, Wentao Zhang, Binhang Yuan, and Conghui He. Efficient pretraining data selection for language models via multi-actor collaboration. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. Association for Computational Linguistics, July 2025.

[4] Yoshua Bengio, Jérôme Louradour, Ronan Collobert, and Jason Weston. Curriculum learning. In *Proceedings of the 26th annual international conference on machine learning*, pages 41–48, 2009.

[5] Jialin Chen, Jan Eric Lenssen, Aosong Feng, Weihua Hu, Matthias Fey, Leandros Tassiulas, Jure Leskovec, and Rex Ying. From similarity to superiority: Channel clustering for time series forecasting. *Advances in Neural Information Processing Systems*, 37:130635–130663, 2024.

[6] Mouxiang Chen, Lefei Shen, Han Fu, Zhuo Li, Jianling Sun, and Chenghao Liu. Calibration of time-series forecasting: Detecting and adapting context-driven distribution shift. In *Proceedings of the 30th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, pages 341–352, 2024.

[7] Si-An Chen, Chun-Liang Li, Sercan O Arik, Nathanael Christian Yoder, and Tomas Pfister. Tsmixer: An all-mlp architecture for time series forecasting. *Transactions on Machine Learning Research*, 2023.

[8] Marco Cuturi and Mathieu Blondel. Soft-dtw: a differentiable loss function for time-series. In *International conference on machine learning*, pages 894–903. PMLR, 2017.

[9] Abhimanyu Das, Weihao Kong, Rajat Sen, and Yichen Zhou. A decoder-only foundation model for time-series forecasting. In *Proceedings of the 41st International Conference on Machine Learning*, 2024.

[10] Daizong Ding, Mi Zhang, Xudong Pan, Min Yang, and Xiangnan He. Modeling extreme events in time series prediction. In *Proceedings of the 25th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining*, pages 1114–1122, 2019.

[11] Wei Fan, Pengyang Wang, Dongkun Wang, Dongjie Wang, Yuanchun Zhou, and Yanjie Fu. Dish-ts: a general paradigm for alleviating distribution shift in time series forecasting. In *Proceedings of the AAAI conference on artificial intelligence*, volume 37, pages 7522–7529, 2023.

[12] Yisong Fu, Fei Wang, Zezhi Shao, Boyu Diao, Lin Wu, Zhulin An, Chengqing Yu, Yujie Li, and Yongjun Xu. On the integration of spatial-temporal knowledge: A lightweight approach to atmospheric time series forecasting, 2025.

[13] Lu Han, Xu-Yang Chen, Han-Jia Ye, and De-Chuan Zhan. Softs: Efficient multivariate time series forecasting with series-core fusion. In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*, 2024.

[14] Lu Han, Han-Jia Ye, and De-Chuan Zhan. Sin: Selective and interpretable normalization for long-term time series forecasting. In *Forty-first International Conference on Machine Learning*, 2024.

[15] Geoffrey Hinton, Oriol Vinyals, and Jeff Dean. Distilling the knowledge in a neural network. *arXiv preprint arXiv:1503.02531*, 2015.

[16] Jincai Huang, Yongjun Xu, Qi Wang, Qi Cheems Wang, Xingxing Liang, Fei Wang, Zhao Zhang, Wei Wei, Boxuan Zhang, Libo Huang, et al. Foundation models and intelligent decision-making: Progress, challenges, and perspectives. *The Innovation*, 2025.

[17] Ming Jin, Shiyu Wang, Lintao Ma, Zhixuan Chu, James Zhang, Xiaoming Shi, Pin-Yu Chen, Yuxuan Liang, Yuan-fang Li, Shirui Pan, et al. Time-llm: Time series forecasting by reprogramming large language models. In *International Conference on Learning Representations*, 2024.

[18] HyunGi Kim, Siwon Kim, Jisoo Mok, and Sungroh Yoon. Battling the non-stationarity in time series forecasting via test-time adaptation. *arXiv preprint arXiv:2501.04970*, 2025.

[19] Taesung Kim, Jinhee Kim, Yunwon Tae, Cheonbok Park, Jang-Ho Choi, and Jaegul Choo. Reversible instance normalization for accurate time-series forecasting against distribution shift. In *International conference on learning representations*, 2021.

[20] Diederik P. Kingma and Jimmy Ba. Adam: A method for stochastic optimization. In *International Conference on Learning Representations*, 2015.

[21] Dilfira Kudrat, Zongxia Xie, Yanru Sun, Tianyu Jia, and Qinghua Hu. Patch-wise structural loss for time series forecasting. *arXiv preprint arXiv:2503.00877*, 2025.

[22] Ying-yee Ava Lau, Zhiwen Shao, and Dit-Yan Yeung. Fast and slow streams for online time series forecasting without information leakage. In *The Thirteenth International Conference on Learning Representations*, 2025.

[23] Vincent Le Guen and Nicolas Thome. Shape and time distortion loss for training deep time series forecasting models. *Advances in neural information processing systems*, 32, 2019.

[24] Hyunwook Lee, Chunggi Lee, Hongkyu Lim, and Sungahn Ko. Tilde-q: a transformation invariant loss function for time-series forecasting. *arXiv preprint arXiv:2210.15050*, 2022.

[25] Chenghan Li, Mingchen Li, and Ruisheng Diao. Tvnet: A novel time series analysis method based on dynamic convolution and 3d-variation. In *The Thirteenth International Conference on Learning Representations*, 2025.

[26] Yaguang Li, Rose Yu, Cyrus Shahabi, and Yan Liu. Diffusion convolutional recurrent neural network: Data-driven traffic forecasting. In *International Conference on Learning Representations*, 2018.

[27] Zhenghao Lin, Zhibin Gou, Yeyun Gong, Xiao Liu, Ruochen Xu, Chen Lin, Yujiu Yang, Jian Jiao, Nan Duan, Weizhu Chen, et al. Not all tokens are what you need for pretraining. *Advances in Neural Information Processing Systems*, 37:29029–29063, 2024.

[28] Haoxin Liu, Harshavardhan Kamarthi, Lingkai Kong, Zhiyuan Zhao, Chao Zhang, and B Aditya Prakash. Time-series forecasting for out-of-distribution generalization using invariant learning. In *Proceedings of the 41st International Conference on Machine Learning*, pages 31312–31325, 2024.

[29] Yong Liu, Tengge Hu, Haoran Zhang, Haixu Wu, Shiyu Wang, Lintao Ma, and Mingsheng Long. itransformer: Inverted transformers are effective for time series forecasting. In *The Twelfth International Conference on Learning Representations*, 2024.

[30] Yong Liu, Haoran Zhang, Chenyu Li, Xiangdong Huang, Jianmin Wang, and Mingsheng Long. Timer: generative pre-trained transformers are large time series models. In *Proceedings of the 41st International Conference on Machine Learning*, pages 32369–32399, 2024.

[31] Jiecheng Lu, Xu Han, Yan Sun, and Shihao Yang. Cats: Enhancing multivariate time series forecasting by constructing auxiliary time series as exogenous variables. In *International Conference on Machine Learning*, pages 32990–33006. PMLR, 2024.

[32] Donghao Luo and Xue Wang. Moderntcn: A modern pure convolution structure for general time series analysis. In *The twelfth international conference on learning representations*, pages 1–43, 2024.

[33] Ziqing Ma, Wenwei Wang, Tian Zhou, Chao Chen, Bingqing Peng, Liang Sun, and Rong Jin. Fusionsf: Fuse heterogeneous modalities in a vector quantized framework for robust solar power forecasting. In *Proceedings of the 30th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, pages 5532–5543, 2024.

[34] Sören Mindermann, Jan M Brauner, Muhammed T Razzak, Mrinank Sharma, Andreas Kirsch, Winnie Xu, Benedikt Höltgen, Aidan N Gomez, Adrien Morisot, Sebastian Farquhar, et al. Prioritized training on points that are learnable, worth learning, and not yet learnt. In *International Conference on Machine Learning*, pages 15630–15649. PMLR, 2022.

[35] Yuqi Nie, Nam H Nguyen, Phanwadee Sinthong, and Jayant Kalagnanam. A time series is worth 64 words: Long-term forecasting with transformers. In *The Eleventh International Conference on Learning Representations*, 2023.

[36] Quang Pham, Chenghao Liu, Doyen Sahoo, and Steven Hoi. Learning fast and slow for online time series forecasting. In *The Eleventh International Conference on Learning Representations*, 2023.

[37] Xiangfei Qiu, Jilin Hu, Lekui Zhou, Xingjian Wu, Junyang Du, Buang Zhang, Chenjuan Guo, Aoying Zhou, Christian S. Jensen, Zhenli Sheng, and Bin Yang. Tfb: Towards comprehensive and fair benchmarking of time series forecasting methods. *Proc. VLDB Endow.*, 17:2363–2377, 2024.

[38] Suman Ravuri, Karel Lenc, Matthew Willson, Dmitry Kangin, Remi Lam, Piotr Mirowski, Megan Fitzsimons, Maria Athanassiadou, Sheleem Kashem, Sam Madge, et al. Skilful precipitation nowcasting using deep generative models of radar. *Nature*, 597(7878):672–677, 2021.

[39] CLEVELAND RB. Stl: A seasonal-trend decomposition procedure based on loess. *J Off Stat*, 6:3–73, 1990.

[40] Sebastian Schmidl, Phillip Wenig, and Thorsten Papenbrock. Anomaly detection in time series: a comprehensive evaluation. *Proceedings of the VLDB Endowment*, 15(9):1779–1797, 2022.

[41] Zezhi Shao, Yujie Li, Fei Wang, Chengqing Yu, Yisong Fu, Tangwen Qian, Bin Xu, Boyu Diao, Yongjun Xu, and Xueqi Cheng. Blast: Balanced sampling time series corpus for universal forecasting models. In *Proceedings of the 31st ACM SIGKDD Conference on Knowledge Discovery and Data Mining V. 2*, pages 2502–2513, 2025.

[42] Zezhi Shao, Fei Wang, Yongjun Xu, Wei Wei, Chengqing Yu, Zhao Zhang, Di Yao, Tao Sun, Guangyin Jin, Xin Cao, et al. Exploring progress in multivariate time series forecasting: Comprehensive benchmarking and heterogeneity analysis. *IEEE Transactions on Knowledge and Data Engineering*, 2024.

[43] Zezhi Shao, Fei Wang, Zhao Zhang, Yuchen Fang, Guangyin Jin, and Yongjun Xu. Hutformer: Hierarchical u-net transformer for long-term traffic forecasting. *arXiv preprint arXiv:2307.14596*, 2023.

[44] Zezhi Shao, Zhao Zhang, Fei Wang, Wei Wei, and Yongjun Xu. Spatial-temporal identity: A simple yet effective baseline for multivariate time series forecasting. In *Proceedings of the 31st ACM international conference on information & knowledge management*, pages 4454–4458, 2022.

[45] Zezhi Shao, Zhao Zhang, Wei Wei, Fei Wang, Yongjun Xu, Xin Cao, and Christian S Jensen. Decoupled dynamic spatial-temporal graph neural network for traffic forecasting. *Proceedings of the VLDB Endowment*, 15(11):2733–2746, 2022.

[46] Xiaoming Shi, Shiyu Wang, Yuqi Nie, Dianqi Li, Zhou Ye, Qingsong Wen, and Ming Jin. Time-moe: Billion-scale time series foundation models with mixture of experts. *arXiv preprint arXiv:2409.16040*, 2024.

[47] Yumin Suh, Bohyung Han, Wonsik Kim, and Kyoung Mu Lee. Stochastic class-based hard example mining for deep metric learning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, June 2019.

[48] Shivakanth Sujit, Somjit Nath, Pedro Braga, and Samira Ebrahimi Kahou. Prioritizing samples in reinforcement learning with reducible loss. *Advances in Neural Information Processing Systems*, 36:23237–23258, 2023.

[49] Tristan Thrush, Christopher Potts, and Tatsunori Hashimoto. Improving pretraining data using perplexity correlations. In *The Thirteenth International Conference on Learning Representations*, 2025.

[50] Ruilin Tong, Yuhang Liu, Javen Qinfeng Shi, and Dong Gong. Coreset selection via reducible loss in continual learning. In *The Thirteenth International Conference on Learning Representations*, 2025.

[51] Hao Wang, Lichen Pan, Yuan Shen, Zhichao Chen, Degui Yang, Yifei Yang, Sen Zhang, Xinggao Liu, Haoxuan Li, and Dacheng Tao. Fredf: Learning to forecast in the frequency domain. In *The Thirteenth International Conference on Learning Representations*, 2025.

[52] Qi Wang, Yiqin Lv, Yixiu Mao, Yun Qu, Yi Xu, and Xiangyang Ji. Robust fast adaptation from adversarially explicit task distribution generation. In *Proceedings of the 31st ACM SIGKDD Conference on Knowledge Discovery and Data Mining V. 1*, pages 1481–1491, 2025.

[53] Qi Wang and Herke Van Hoof. Doubly stochastic variational inference for neural processes with hierarchical latent variables. In *International Conference on Machine Learning*, pages 10018–10028. PMLR, 2020.

[54] Qi Cheems Wang, Zehao Xiao, Yixiu Mao, Yun Qu, Jiayi Shen, Yiqin Lv, and Xiangyang Ji. Model predictive task sampling for efficient and robust adaptation. *arXiv preprint arXiv:2501.11039*, 2025.

[55] Shiyu Wang, Haixu Wu, Xiaoming Shi, Tengge Hu, Huakun Luo, Lintao Ma, James Y Zhang, and JUN ZHOU. Timemixer: Decomposable multiscale mixing for time series forecasting. In *The Twelfth International Conference on Learning Representations*, 2024].

[56] Yuxuan Wang, Haixu Wu, Jiaxiang Dong, Guo Qin, Haoran Zhang, Yong Liu, Yunzhong Qiu, Jianmin Wang, and Mingsheng Long. Timexer: Empowering transformers for time series forecasting with exogenous variables. In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*, 2024.

[57] Zhiyuan Wang, Xovee Xu, Goce Trajcevski, Kunpeng Zhang, Ting Zhong, and Fan Zhou. Pref: Probabilistic electricity forecasting via copula-augmented state space model. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 36, pages 12200–12207, 2022.

[58] Andreas S Weigend, Morgan Mangeas, and Ashok N Srivastava. Nonlinear gated experts for time series: Discovering regimes and avoiding overfitting. *International journal of neural systems*, 6(04):373–399, 1995.

[59] Qingsong Wen, Weiqi Chen, Liang Sun, Zhang Zhang, Liang Wang, Rong Jin, Tieniu Tan, et al. Onenet: Enhancing time series forecasting models under concept drift by online ensembling. *Advances in Neural Information Processing Systems*, 36:69949–69980, 2023.

[60] Guillaume Wenzek, Marie-Anne Lachaux, Alexis Conneau, Vishrav Chaudhary, Francisco Guzmán, Armand Joulin, and Edouard Grave. Ccnet: Extracting high quality monolingual datasets from web crawl data, 2019.

[61] Alexander Wettig, Aatmik Gupta, Saumya Malik, and Danqi Chen. Qurating: Selecting high-quality data for training language models. In *Forty-first International Conference on Machine Learning*. PMLR, 2024.

[62] Gerald Woo, Chenghao Liu, Akshat Kumar, Caiming Xiong, Silvio Savarese, and Doyen Sahoo. Unified training of universal time series forecasting transformers. In *Proceedings of the 41st International Conference on Machine Learning*, pages 53140–53164, 2024.

[63] Haixu Wu, Tengge Hu, Yong Liu, Hang Zhou, Jianmin Wang, and Mingsheng Long. Timesnet: Temporal 2d-variation modeling for general time series analysis. In *The Eleventh International Conference on Learning Representations*, 2023.

[64] Haixu Wu, Jiehui Xu, Jianmin Wang, and Mingsheng Long. Autoformer: Decomposition transformers with auto-correlation for long-term series forecasting. *Advances in neural information processing systems*, 34:22419–22430, 2021.

[65] Haixu Wu, Hang Zhou, Mingsheng Long, and Jianmin Wang. Interpretable weather forecasting for worldwide stations with a unified deep model. *Nature Machine Intelligence*, 5(6):602–611, 2023.

[66] Xingjian Wu, Xiangfei Qiu, Zhengyu Li, Yihang Wang, Jilin Hu, Chenjuan Guo, Hui Xiong, and Bin Yang. CATCH: Channel-aware multivariate time series anomaly detection via frequency patching. In *The Thirteenth International Conference on Learning Representations*, 2025.

[67] Zonghan Wu, Shirui Pan, Guodong Long, Jing Jiang, Xiaojun Chang, and Chengqi Zhang. Connecting the dots: Multivariate time series forecasting with graph neural networks. In *Proceedings of the 26th ACM SIGKDD international conference on knowledge discovery & data mining*, pages 753–763, 2020.

[68] Sang Michael Xie, Hieu Pham, Xuanyi Dong, Nan Du, Hanxiao Liu, Yifeng Lu, Percy Liang, Quoc V. Le, Tengyu Ma, and Adams Wei Yu. Doremi: Optimizing data mixtures speeds up language model pretraining. *Advances in Neural Information Processing Systems*, 36:69798–69818, 2023.

[69] Sang Michael Xie, Shibani Santurkar, Tengyu Ma, and Percy S Liang. Data selection for language models via importance resampling. *Advances in Neural Information Processing Systems*, 36:34201–34227, 2023.

[70] Zhijian Xu, Ailing Zeng, and Qiang Xu. Fits: Modeling time series with $10k$ parameters. In *The Twelfth International Conference on Learning Representations*, 2024.

[71] Yiyuan Yang, Chaoli Zhang, Tian Zhou, Qingsong Wen, and Liang Sun. Dcdetector: Dual attention contrastive representation learning for time series anomaly detection. In *Proceedings of the 29th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, pages 3033–3045, 2023.

[72] Chengqing Yu, Fei Wang, Zezhi Shao, Tangwen Qian, Zhao Zhang, Wei Wei, Zhulin An, Qi Wang, and Yongjun Xu. Ginar+: A robust end-to-end framework for multivariate time series forecasting with missing values. *IEEE Transactions on Knowledge and Data Engineering*, 37(8):4635–4648, 2025.

[73] Chengqing Yu, Fei Wang, Zezhi Shao, Tangwen Qian, Zhao Zhang, Wei Wei, and Yongjun Xu. Ginar: An end-to-end multivariate time series forecasting model suitable for variable missing. In *Proceedings of the 30th ACM SIGKDD conference on knowledge discovery and data mining*, pages 3989–4000, 2024.

[74] Chengqing Yu, Fei Wang, Zezhi Shao, Tao Sun, Lin Wu, and Yongjun Xu. Dsformer: A double sampling transformer for multivariate time series long-term prediction. In *Proceedings of the 32nd ACM international conference on information and knowledge management*, pages 3062–3072, 2023.

[75] Chengqing Yu, Fei Wang, Chuanguang Yang, Zezhi Shao, Tao Sun, Tangwen Qian, Wei Wei, Zhulin An, and Yongjun Xu. Merlin: Multi-view representation learning for robust multivariate time series forecasting with unfixed missing rates. In *Proceedings of the 31st ACM SIGKDD Conference on Knowledge Discovery and Data Mining V. 2*, pages 3633–3644, 2025.

[76] Ailing Zeng, Muxi Chen, Lei Zhang, and Qiang Xu. Are transformers effective for time series forecasting? In *Proceedings of the AAAI conference on artificial intelligence*, volume 37, pages 11121–11128, 2023.

[77] Chi Zhang, Huaping Zhong, Kuan Zhang, Chengliang Chai, Rui Wang, Xinlin Zhuang, Tianyi Bai, Qiu Jiantao, Lei Cao, Ju Fan, Ye Yuan, Guoren Wang, and Conghui He. Harnessing diversity for important data selection in pretraining large language models. In *The Thirteenth International Conference on Learning Representations*, 2025.

[78] Yuchen Zhang, Mingsheng Long, Kaiyuan Chen, Lanxiang Xing, Ronghua Jin, Michael I Jordan, and Jianmin Wang. Skilful nowcasting of extreme precipitation with nowcastnet. *Nature*, 619(7970):526–532, 2023.

[79] Yunhao Zhang and Junchi Yan. Crossformer: Transformer utilizing cross-dimension dependency for multivariate time series forecasting. In *The eleventh international conference on learning representations*, 2023.

[80] Tianjie Zhao, Sheng Wang, Chaojun Ouyang, Min Chen, Chenying Liu, Jin Zhang, Long Yu, Fei Wang, Yong Xie, Jun Li, et al. Artificial intelligence for geoscience: Progress, challenges, and perspectives. *The Innovation*, 5(5), 2024.

[81] Haoyi Zhou, Shanghang Zhang, Jieqi Peng, Shuai Zhang, Jianxin Li, Hui Xiong, and Wancai Zhang. Informer: Beyond efficient transformer for long sequence time-series forecasting. In *Proceedings of the AAAI conference on artificial intelligence*, volume 35, pages 11106–11115, 2021.

[82] Tian Zhou, Ziqing Ma, Qingsong Wen, Xue Wang, Liang Sun, and Rong Jin. Fedformer: Frequency enhanced decomposed transformer for long-term series forecasting. In *International conference on machine learning*, pages 27268–27286. PMLR, 2022.

[83] Tian Zhou, Peisong Niu, Liang Sun, Rong Jin, et al. One fits all: Power general time series analysis by pretrained lm. *Advances in neural information processing systems*, 36:43322–43355, 2023.

A Proofs

A.1 Assumptions

Assumption 2 (Lipschitz Continuity). *We assume that f reserves the Lipschitz continuity w.r.t. $\boldsymbol{\theta}$, i.e., $\forall \boldsymbol{\theta}_1, \boldsymbol{\theta}_2 \in \boldsymbol{\Theta}$ satisfying*

$$||f(\mathbf{X}; \boldsymbol{\theta}_1) - f(\mathbf{X}; \boldsymbol{\theta}_2)|| \le L_f ||\boldsymbol{\theta}_1 - \boldsymbol{\theta}_2||, \tag{15}$$

where L_f is the Lipschitz constant.

Justification It can be ensured by the Lipschitz-continuous activations of the neural network and the continuously differentiable MSE loss.

Assumption 3 (Bounded Prediction Residual). *We assume that there exists a constant $R > 0$ such that*

$$|\epsilon_t| \le R, \quad \forall t \in \{1, \dots, T\}. \tag{16}$$

Justification Empirically, residuals often follow a light-tailed distribution, where extreme deviations are rare. Therefore, there exists a finite high-probability bound.

Assumption 4 (Bounded Gradient). *We assume that there exists a constant $G > 0$ such that*

$$||\nabla_{\boldsymbol{\theta}_\tau} \mathcal{L}|| < G, \quad \forall \tau \in \mathbb{Z}^+. \tag{17}$$

Justification The boundedness of gradients is ensured by Lipschitz-continuous activations and weight constraints of the neural network, preventing explosive updates and ensuring stable optimization. Gradient clipping can further enforce it.

A.2 Proof of Theorem 1

Proof. Let τ_i denote the training time corresponding to the i-th most recent prediction residual $\epsilon_t^{(i)}$, such that $\epsilon_t(\boldsymbol{\theta}_{\tau_i}) \equiv \epsilon_t^{(i)}$.

$$
\begin{aligned}
|\hat{\sigma}_t^2 - \hat{\sigma}_t^2(\boldsymbol{\theta}_\tau)| &\le \frac{1}{n_t} \sum_{i=1}^{n_t} |\epsilon_t^2(\boldsymbol{\theta}_{\tau_i}) - \epsilon_t^2(\boldsymbol{\theta}_\tau) - n_t(\bar{\epsilon}_t^2 - \bar{\epsilon}_t^2(\boldsymbol{\theta}_\tau))| \\
&\le \frac{1}{n_t} \sum_{i=1}^{n_t} |\epsilon_t^2(\boldsymbol{\theta}_{\tau_i}) - \epsilon_t^2(\boldsymbol{\theta}_\tau)| + |\bar{\epsilon}_t^2 - \bar{\epsilon}_t^2(\boldsymbol{\theta}_\tau)| \\
&\le 2 \max_i |\epsilon_t^2(\boldsymbol{\theta}_{\tau_i}) - \epsilon_t^2(\boldsymbol{\theta}_\tau)| \\
&\le 2 \max_i |\epsilon_t(\boldsymbol{\theta}_{\tau_i}) - \epsilon_t(\boldsymbol{\theta}_\tau)| \cdot |\epsilon_t(\boldsymbol{\theta}_{\tau_i}) + \epsilon_t(\boldsymbol{\theta}_\tau)| \\
&= 2 \max_i |f(\mathbf{X}, \boldsymbol{\theta}_{\tau_i})_t - f(\mathbf{X}, \boldsymbol{\theta}_\tau)_t| \cdot |\epsilon_t(\boldsymbol{\theta}_{\tau_i}) + \epsilon_t(\boldsymbol{\theta}_\tau)| \\
&\le 2 \max_i L_f ||\boldsymbol{\theta}_{\tau_i} - \boldsymbol{\theta}_\tau|| \cdot 2R \quad \text{(Assumption 2 and 3)} \\
&\le 4 L_f R \eta G \max_i |\tau - \tau_i|. \quad \text{(Assumption 4)}
\end{aligned}
\tag{18}
$$

Given that adjacent epochs are separated by no more than $2K - 1$ iterations, we obtain:

$$|\hat{\sigma}_t^2 - \hat{\sigma}_t^2(\boldsymbol{\theta}_\tau)| \le 4 L_f R \eta G(2K - 1). \tag{19}$$

$\square$

B   Related Work

B.1   Data Selection Strategies in Deep Learning

Data selection strategies aim to improve the efficiency, generalizability, and robustness of deep learning models by carefully selecting subsets of data for training. Instead of using the entire dataset indiscriminately, these methods prioritize samples that contribute more significantly to model learning dynamics or downstream performance. Some studies such as curriculum learning [4] and hard sample mining [47], select or filter samples based on their *difficulties*. In LM pretraining, samples or tokens are typically selected for high *quality* [68, 61, 27, 3], great *importance* [69], or strong *diversity* [77] to enhance both training efficiency and generalization. Data filtering pipelines often rely on heuristic metrics, such as perplexity [60, 49] and toxicity, or learned metrics [34, 27, 48], to remove duplicated, or low-quality data. Unlike the above studies, our approach is closely tied to the characteristics of time series data, focusing on filtering noise or anomalies. By identifying and removing such irregular patterns, the model can learn from representative and generalizable timesteps and achieve more stable and reliable predictions.

C   Implementation Details

C.1   Dataset Descriptions

We conduct experiments on 8 real-world datasets to evaluate the effectiveness of the proposed selective learning, including:

- **ETT** (Electricity Transformer Temperature) [81] contains 7 features of electricity transformer data collected from two separate counties from July 2016 to July 2018. It contains four datasets: **ETTh1, ETTh2, ETTm1, ETTm2**, where ETTh1 and ETTh2 are recorded every hour, and ETTm1 and ETTm2 are recorded every 15 minutes.

- **Electricity** [64] records the hourly electricity consumption data of 321 clients from 2012 to 2014. Each variable represents a client's electricity consumption.

- **Exchange** [64] collects the panel data of daily exchange rates from 1990 to 2016 from 8 countries, including Australia, Britain, Canada, Switzerland, China, Japan, New Zealand, and Singapore.

- **Weather**[64] includes 21 meteorological factors collected every 10 minutes from the weather station of the Max Planck Biogeochemistry Institute in 2020.

- **ILI** (Influenza-Like Illness) [64] includes the weekly recorded patient data from the Centers for Disease Control and Prevention of the United States between 2002 and 2021.

We follow the same data processing and train-validation-test set split protocol used in TimesNet[63], where the train, validation, and test datasets are strictly divided according to chronological order to ensure no data leakage issues. The statistics of the datasets are provided in Table 5.

Table 5: Statistics of the datasets.

Dataset	Dim	Prediction Length	Dataset Size	Split	Frequency	Domain
ETTh1, ETTh2	7	{96, 192, 336, 720}	14,400	6:2:2	Hourly	Electricity
ETTm1, ETTm2	7	{96, 192, 336, 720}	57,600	6:2:2	15min	Electricity
Exchange	8	{96, 192, 336, 720}	7,588	7:1:2	Daily	Economy
Weather	21	{96, 192, 336, 720}	52,696	7:1:2	10min	Weather
Electricity	321	{96, 192, 336, 720}	26,304	7:1:2	Hourly	Electricity
ILI	7	{24, 36, 48, 60}	966	7:1:2	Weekly	Health

C.2 Baselines

- **Informer** [81] is a Transformer for time series forecasting (TSF) with a ProbSparse self-attention mechanism.
- **Crossformer** [79] utilizes attention to capture both temporal and multivariate correlations.
- **PatchTST** [35] splits the input time series into patches, which serve as input tokens of the Transformer. It proposes a channel-independent strategy.
- **TimesNet** [63] transforms time series into 2D tensors and employs CNN to capture inter- and intra-period dependencies.
- **iTransformer** [29] embeds each series independently to the variate token and applies self-attention to capture multivariate correlations.
- **TimeMixer** [55] is an MLP-based model that captures multi-scale patterns by decomposing time series into different scales and mixing them through MLP layers.

Notably, PatchTST, iTransformer, and TimeMixer are equipped with RevIN [19] to handle the distribution shift issue. Selective learning can provide additional performance gains while maintaining full compatibility with existing normalization techniques.

Algorithm 1 The workflow of selective learning.

1: **INPUT:** The model $f(\cdot, \boldsymbol{\theta})$, the training set $\mathcal{D}_{train} = \{(\mathbf{X}_{t-L:t}, \mathbf{X}_{t:t+F})\}_{t=L}^{T-F}$, the estimation model $g(\cdot, \boldsymbol{\phi})$ trained on $\mathcal{D}_{train}$, and the number of iterations N_{it}.
2: **OUTPUT:** Optimized model $f(\cdot, \boldsymbol{\theta}_\tau)$.
3: Initialize $f(\cdot, \boldsymbol{\theta}_0)$ and a historical residual archive S
4: **for** τ *in* $\{0, 1, \cdot, N_{it} - 1\}$ **do**
5:    $\hat{\mathbf{X}}_{t:t+F} = f(\mathbf{X}_{t-L:t}; \boldsymbol{\theta}_\tau)$ // Forward
6:    // Calculate the residual and update S
7:    $\boldsymbol{\epsilon}_{t:t+F} = \mathbf{X}_{t:t+F} - \hat{\mathbf{X}}_{t:t+F}$
8:    $S \leftarrow \boldsymbol{\epsilon}_{t:t+F}$
9:    // Uncertainty mask
10:    **if** $\tau \geq K$ **then**
11:      // Update the residual entropy once per epoch
12:      **if** $\tau \% K = 0$ **then**
13:        **for** t *in* $\{0, \cdots, T\}$ **do**
14:          Calculate $\hat{H}(\epsilon_t)$ by Eq.(8)
15:        **end for**
16:        $\gamma_u = \text{Top-}r_u\% \ H(\epsilon_t)$ *for* $t \in \{0, \cdots, T-1\}$
17:      **end if**
18:      Calculate $\mathcal{M}_u^{(\tau)}$ by Eq.(11)
19:    **end if**
20:    // Anomaly Mask
21:    Calculate $\hat{S}(X_t)$ by Eq.(13)
22:    $\gamma_a = \text{Top-}r_a\% \ S(X_t)$ *for* $t \in \{0, \cdots, F-1\}$
23:    Calculate $\mathcal{M}_a^{(\tau)}$ by Eq.(14)
24:    $\mathcal{M}^{(\tau)} = \mathcal{M}_u^{(\tau)} \vee \mathcal{M}_a^{(\tau)}$
25:    $\mathcal{L}_{SL} = \frac{1}{N \cdot |\mathcal{M}^{(\tau)}|} \sum_{i=0}^{F-1} ||\mathcal{M}^{(\tau)}(X_{t+i} - f(\mathbf{X}_{t-L:t}; \boldsymbol{\theta}_\tau)_i)||^2$
26:    $\boldsymbol{\theta}_{\tau+1} = \boldsymbol{\theta}_\tau - \eta \nabla_{\boldsymbol{\theta}} \mathcal{L}_{SL}$ // Optimization through selective learning loss
27: **end for**

C.3 Implementation Details

In this section, we provide the implementation details of selective learning. The overall workflow of selective learning is provided in Algorithm 1.

Uncertainty Mask   We employ a global threshold γ_u computed across the entire training sequence $\mathbf{X}_{0:T}$ to determine uncertainty masks. This choice prevents performance degradation caused by

masking normal (but relatively uncertain within specific samples) timesteps when masking ratios increase. For computational efficiency, the threshold is updated only once per epoch. This design additionally mitigates the cold-start issue of uncertainty masking in the first epoch, where inadequate residual entropy estimates would otherwise lead to suboptimal masking decisions.

When processing extremely large datasets that exceed memory capacity for saving full-sequence residuals, we recommend two alternative approaches: (1) adopting per-sample thresholds with conservatively low masking ratios, or (2) implementing a max-heap algorithm for threshold computation. Both solutions maintain computational feasibility while preserving the benefits of uncertain masking.

Anomaly Mask We employ a per-sample threshold γ_a computed across the prediction sequence $\mathbf{X}_{t:t+F}$ to determine uncertainty masks. Global anomaly masking over $\mathbf{X}_{0:T}$ is adversely affected by training dynamics (where later samples consistently produce smaller residuals), resulting in suboptimal mask selection.

In this paper, the estimation model g is trained on the entire training set of f, diverging from some existing works in other fields that employ held-out sets to train auxiliary models to prevent overfitting in the training set [34]. We prevent overfitting instead by employing a simple model (e.g., DLinear) as g, thereby safeguarding the main model's performance. Additionally, unlike other data modalities, the patterns in time-series data often change over time. Consequently, using a holdout set may lead to underfitting of certain patterns in g, particularly for datasets with strong non-stationarity. When using a simple model as g, it should be trained until full convergence on the training set to avoid underfitting. Additionaly, since the sample partitioning in TSF depends on the sliding window size, each forecasting window size necessitates training a distinct model.

C.4 Selection of Masking Ratio

The masking ratios are crucial hyperparameters in selective learning. We provide the default masking ratio in Table 6. However, the optimal masking ratio may vary across different models. Users can perform a hyperparameter search around the default masking rate to find the optimal masking ratio. For a new dataset, we recommend a three-stage optimization protocol: (1) optimize the noise masking ratio to achieve peak performance; (2) fix the noise masking ratio and tune the anomaly masking ratio to its optimal value; (3) conduct a local hyperparameter search within the neighborhood of these determined masking ratios.

Additionally, specific masking ratios may also be attempted as initial values:

- $r_u = 10\%, r_a = 10\%$ for stable and high-quality datasets.
- $r_u = 30\%, r_a = 30\%$ for datasets containing certain levels of noise and anomalies.
- $r_a = 90\%$ for highly-volatile dataset.

Table 6: The default masking ratio for the datasets.

Dataset	ETTh1	ETTh2	ETTm1	ETTm2	Electricity	Exchange	Weather	ILI
Uncertainty Mask	30%	10%	20%	20%	10%	/	10%	10%
Anomaly Mask	30%	60%	20%	50%	10%	90%	20%	10%

D Running Cost

In this section, we analyze the computational cost introduced by selective learning from the perspective of running time and memory usage.

Complexity Analysis We first theoretically analyze of the time and space complexity of selective learning algorithm.

- **Time Complexity**: Let B be the batch size. For uncertainty mask, residual entropy updates cost $\mathcal{O}(BFN)$ per epoch. The complexity of the anomaly mask depends on the architecture of the estimation model. Taking a linear model as an example, the forward pass of the

estimation model requires $\mathcal{O}(BLN)$ complexity, and the masking process adds $\mathcal{O}(BFN)$. Therefore, the complexity of the anomaly mask is $\mathcal{O}(B(L+F)N)$.

- **Space Complexity**: Storing residuals for T timesteps requires $\mathcal{O}(TFN)$ space.

Running Time We measure the running time per epoch of different models trained without and with selective learning on the ETTh1 and ETTm2 datasets. The results in Table 7 demonstrate that selective learning maintains computational efficiency, adding acceptable training time while achieving significant performance gains.

Table 7: Running cost of selective learning (SL). The results are averaged over 3 runs.

Method		TimesNet				iTransformer				TimeMixer		
Metric	w/o (s/Epoch)	SL (s/Epoch)	Time Inc. (s/Epoch)	MSE Dec. %	w/o (s/Epoch)	SL (s/Epoch)	Time Inc. (s/Epoch)	MSE Dec. %	w/o (s/Epoch)	SL (s/Epoch)	Time Inc. (s/Epoch)	MSE Dec. %
ETTh1 96	7.2	9.0	1.8	10.6%	1.8	2.0	0.2	7.7%	5.0	5.4	0.4	6.9%
192	10.2	11.0	0.8	13.0%	1.9	2.2	0.3	7.0%	5.1	5.6	0.5	6.8%
336	10.2	12.4	2.2	12.9%	2.1	2.3	0.2	8.1%	5.1	5.8	0.7	4.0%
720	14.1	15.7	1.6	18.9%	2.2	3.1	0.9	13.6%	5.3	6.5	1.2	10.5%
Avg.	10.4	12.0	1.6	14.0%	2.0	2.4	0.4	9.3%	5.1	5.8	0.7	6.9%
ETTm2 96	43.2	54.5	11.3	11.5%	9.7	11.0	1.3	5.1%	18.1	18.8	0.7	4.7%
192	50.1	56.6	6.5	8.1%	10.2	11.2	1.0	9.0%	19.2	20.0	0.8	5.6%
336	56.3	64.2	7.9	13.1%	10.9	11.8	0.9	7.2%	20.5	21.1	0.6	5.7%
720	77.4	100.5	23.1	10.0%	10.3	14.2	3.9	4.2%	19.1	24.0	4.9	4.3%
Avg.	56.8	69.0	12.2	10.7%	10.3	12.0	2.7	6.4%	19.2	21.0	1.8	5.0%

Memory Usage Our implementation processes historical residuals on the CPU, resulting in merely 2MB of additional GPU memory allocation. Although migrating these operations to the GPU would improve computational throughput, it would increase GPU memory consumption. Additionally, maintaining historical residuals in memory requires approximately $4|\mathcal{D}_{train}|NF$ bytes of RAM (e.g., <0.1GB RAM for ETTh1 and about 7GB for Electricity when $F = 336$).

E Discussion

E.1 Non-generalizable Timesteps vs. Distribution Shift

In recent years, distribution shift in non-stationary time series has been widely studied by the research community [19, 11, 14] and shares conceptual similarities with non-generalizable timestemps. This section compares and contrasts non-generalizable time steps with distribution shift. At their core, both concepts describe a fundamental challenge in deep time series forecasting: the problem of a mismatch between the data a model was trained on and the data it encounters in test, which is typically induced by changes in environmental or exogenous variables

However, the crucial distinction lies in their scale: distribution shift is typically a instance- or segment-level phenomenon, where the statistics of the entire dataset change gradually or abruptly over time. In contrast, a non-generalizable timestep is often a localized, point-level issue. It refers to an individual or a small set of timesteps whose patterns are uncertain or anomalous that they cannot be reliably learned or predicted by the model, even if the overall data distribution remains stable.

Therefore, our proposed selective learning offers a finer-grained solution to prevent models from being affected by non-generalizable data.

E.2 Dynamic Masking vs. Static Masking

Static masking offers an intuitive implementation approach, for example, training an estimation model to estimate the distribution of each timestep, or employing time series anomaly detection models [40, 71, 66] to target anomalous timesteps. In contrast, dynamic masking adaptively modifies the masked timesteps during training. This approach offers two key advantages:

- Static masking significantly alters the distribution of $\mathcal{D}_{train}$, thereby introducing bias, whereas dynamic masking adapts the mask during training to mitigate this in expectation. Specifically, a given timestep may be masked only during certain training phases and within particular lookback windows, while remaining in other contexts and stages.

- For rare but critical extreme events (e.g., extreme weather)[10, 78] that are less generalizable, dynamic masking first learns the most generalizable timesteps and gradually attempts to learn timesteps previously considered anomalies. Static masking, by contrast, consistently excludes these patterns, resulting in compromised forecasting capacity for extreme events.

We present a comparison between dynamic and static masking in Table 8 using iTransformer, showing that dynamic masking yields consistently better performance, which validates our claims.

Table 8: Comparison of static and dynamic masking.

Method	Metric	ETTh1	ETTm2	Exchange
Static masking	MSE	0.426	0.264	0.358
	MAE	0.437	0.324	0.408
Dynamic masking	MSE	**0.415**	**0.256**	**0.343**
	MAE	**0.425**	**0.313**	**0.399**

E.3 Capacity of Handling Clean Datasets

If selective learning is employed, then even clean datasets will be masked. This section conducts additional experiments and make discussions to investigate whether selective learning has an impact on model performance on clean datasets.

To evaluate selective learning under clean conditions, we construct a synthetic dataset that is theoretically clean (without any noise or anomalies). An ideal time series without noise or anomalies can be decomposed into trend and periodic components [39, 64]. Accordingly, we synthesize the dataset by combining a linear trend components with daily, weekly, and yearly sinusoidal patterns to generate multivariate time series. For each channel, both the trend slope and the amplitudes of each periodic component are sampled uniformly from specified ranges.

We have conducted experiments on the synthetic dataset using iTransformer as the backbone model. As shown in Table 9, uncertainty masking can reduce model performance on clean datasets (by misidentifying noise), while anomaly masking maintains or even marginally improves performance. This effect stems from our dynamic masking approach for anomalies: When a normal timestep is mistakenly masked in one epoch, it can likely be learned in subsequent epochs.

In summary, we recommend using the anomaly mask on clean datasets, while exercising caution when applying uncertainty masking. Given that real-world time series datasets typically contain a certain level of noise and anomalies, it is therefore beneficial to select an appropriate masking ratio.

Table 9: iTransformer's performance on the synthetic dataset ($L = 336, F = 336$). Better results with selective learning are in **bold**.

Method	iTransformer	+ Uncertainty mask			+Anomaly mask		
		r_u=5%	r_u=10%	r_u=20%	r_a=5%	r_a=10%	r_a=20%
MSE	0.0295	0.0475	0.1569	0.1640	0.0299	**0.0294**	**0.0293**
MAE	0.0838	**0.0791**	0.0848	0.1004	**0.0823**	**0.0829**	**0.0833**

F   Limitations and Future Work

Beyond Forecasting Task   Our work currently focuses on time series forecasting tasks. However, our idea can also be applied to other time series analysis tasks, such as imputation and classification [53], guiding the model to focus more on generalizable patterns. While we highlight these potential extensions as promising directions, a thorough investigation of their applicability and effectiveness remains an open question. We leave this exploration for future work.

Extreme Event Forecasting Capacity   The dual-masking mechanism may filter out rare extreme events present in the training set. Although dynamic masking can mitigate this effect, the model's

predictive capability for extreme events may still be compromised. For scenarios where extreme events are critically important, we recommend fine-tuning the selective learning-trained model using online learning [59, 36, 22] or test-time adaptation [2, 6, 18, 52] after deployment. For example, SOLID [6] retrieves training samples similar to the current input (including potentially masked ones) to fine-tune the prediction head.

Pretraining for Time Series Foundation Model Recently, time series foundation models (TSFMs) have achieved rapid advancements [62, 9, 30, 1, 46, 41, 16]. Selective learning currently focuses on in-domain forecasting. We leave this as future work. Since the TSFM is trained on samples drawn from a large-scale dataset, we cannot estimate the residual entropy and lower bound at each timestep. As a result, the existing design is not directly compatible with the training of TSFMs. However, thanks to the strong representational capacity of TSFMs, we can train a probabilistic forecasting model (e.g., Chronos [1], MOIRAI [62]) to directly predict the distribution of each timestep, thereby selecting the generalizable timesteps. Future work can further combine the selective learning with some model predictive methods [54] for TSFMs' robustness and efficiency improvement.

G Full Experimental Results

G.1 Full Forecasting Results

The full forecasting results are provided in Table 10 and 11 due to the page limitation of the main text. It can be observed that selective learning significantly improves models' performance in all cases, demonstrating its effectiveness.

G.2 Full Results of Training Objective Comparison

The full results of the training objective comparison are provided in Table 12 due to the page limitation of the main text. It is evident that selective learning consistently achieves superior performance. These results demonstrate that global alignment over whole sequences, whether in shape or distribution, in the temporal or frequency domain, proves suboptimal, validating the effectiveness of selective learning.

G.3 Full Ablation Results

The full ablation results are provided in Table 13. The results demonstrate that the model with full selective learning consistently achieves the best performance. Removing either mask leads to significant performance degradation across all four datasets. Additionally, replacing the dual-mask mechanism with random masking reduces model performance to levels comparable to or worse than the unmasked counterparts. This suggests that the effectiveness of selective learning fundamentally stems from our dual-mask mechanism, which directs model attention to *generalizable* timesteps while filtering out *non-generalizable* ones.

H Case Study

To showcase the effectiveness of selective learning, we provide supplementary prediction cases of five baselines across five representative datasets in Figure 6. The visualizations clearly show that selective learning can enhance models' forecasting performance and generalizability.

Table 10: Part 1 of the full forecasting results without/with selective learning (SL). Better results are in **bold**. Avg. denotes the average from all prediction lengths, and Δ denotes the averaged improvements caused by selective learning.

Method		Informer				Crossformer				PatchTST			
		w/o		+SL		w/o		+SL		w/o		+SL	
Metric		MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE
ETTh1	96	1.087	0.831	**0.478**	**0.487**	0.406	0.426	**0.373**	**0.397**	0.377	0.397	**0.368**	**0.386**
	192	1.262	0.917	**0.492**	**0.490**	0.448	0.455	**0.417**	**0.422**	0.417	0.421	**0.412**	**0.413**
	336	1.379	0.952	**0.503**	**0.526**	0.460	0.466	**0.448**	**0.443**	0.448	0.442	**0.433**	**0.426**
	720	1.428	0.968	**0.680**	**0.634**	0.505	0.512	**0.484**	**0.500**	0.465	0.472	**0.425**	**0.443**
	Avg.	1.289	0.917	**0.538**	**0.534**	0.455	0.465	**0.431**	**0.441**	0.427	0.433	**0.410**	**0.417**
	Δ			-58.3%	-41.8%			-5.33%	-5.22%			-4.10%	-3.70%
ETTh2	96	3.107	1.475	**1.516**	**0.937**	0.717	0.579	**0.474**	**0.463**	0.311	0.367	**0.296**	**0.353**
	192	3.707	1.659	**1.697**	**1.012**	0.736	0.609	**0.610**	**0.548**	0.381	0.410	**0.369**	**0.400**
	336	2.671	1.346	**1.595**	**0.982**	0.739	0.621	**0.618**	**0.553**	0.418	0.437	**0.401**	**0.428**
	720	2.543	1.348	**2.065**	**1.175**	1.113	0.784	**0.935**	**0.707**	0.443	0.464	**0.431**	**0.452**
	Avg.	3.007	1.457	**1.718**	**1.027**	0.826	0.648	**0.659**	**0.568**	0.388	0.420	**0.374**	**0.408**
	Δ			-42.9%	-29.5%			-20.2%	-12.4%			-3.61%	-2.68%
ETTm1	96	0.443	0.446	**0.304**	**0.350**	0.307	0.360	**0.301**	**0.353**	0.294	0.343	**0.290**	**0.333**
	192	0.618	0.573	**0.344**	**0.374**	0.366	0.404	**0.344**	**0.381**	0.339	0.373	**0.333**	**0.359**
	336	0.865	0.701	**0.374**	**0.395**	0.446	0.453	**0.403**	**0.422**	0.372	0.393	**0.367**	**0.381**
	720	0.941	0.757	**0.431**	**0.433**	0.578	0.531	**0.503**	**0.496**	0.424	0.428	**0.420**	**0.414**
	Avg.	0.717	0.619	**0.363**	**0.388**	0.424	0.437	**0.388**	**0.413**	0.357	0.385	**0.353**	**0.372**
	Δ			-49.3%	-37.3%			-8.60%	-5.49%			-1.33%	-3.32%
ETTm2	96	0.334	0.443	**0.226**	**0.345**	0.281	0.356	**0.194**	**0.291**	0.175	0.261	**0.164**	**0.251**
	192	0.729	0.676	**0.367**	**0.464**	0.374	0.450	**0.262**	**0.341**	0.236	0.306	**0.218**	**0.289**
	336	1.416	0.955	**0.595**	**0.610**	0.676	0.582	**0.391**	**0.430**	0.293	0.347	**0.266**	**0.322**
	720	3.460	1.601	**0.786**	**0.690**	1.019	0.724	**0.631**	**0.570**	0.379	0.402	**0.361**	**0.384**
	Avg.	1.485	0.919	**0.494**	**0.527**	0.588	0.528	**0.370**	**0.408**	0.271	0.329	**0.252**	**0.312**
	Δ			-66.8%	-42.6%			-37.1%	-22.7%			-6.83%	-5.32%
Eletricity	96	0.300	0.389	**0.266**	**0.364**	0.137	0.236	**0.134**	**0.228**	0.139	0.235	0.139	**0.234**
	192	0.311	0.399	**0.286**	**0.384**	0.162	0.260	**0.147**	**0.241**	0.153	0.248	**0.152**	**0.245**
	336	0.349	0.434	**0.301**	**0.391**	0.190	0.283	**0.169**	**0.271**	0.168	0.267	**0.166**	**0.258**
	720	0.406	0.459	**0.316**	**0.404**	0.237	0.330	**0.220**	**0.315**	0.208	0.296	**0.204**	**0.293**
	Avg.	0.342	0.420	**0.292**	**0.386**	0.182	0.277	**0.168**	**0.264**	0.167	0.262	**0.165**	**0.258**
	Δ			-14.4%	-8.21%			-7.71%	-4.87%			-1.05%	-1.53%
Exchange	96	0.906	0.763	**0.408**	**0.514**	0.289	0.396	**0.205**	**0.328**	0.078	0.196	0.078	0.196
	192	1.291	0.908	**0.621**	**0.638**	0.527	0.558	**0.355**	**0.447**	0.161	0.290	**0.157**	**0.287**
	336	1.334	0.953	**0.825**	**0.745**	0.858	0.719	**0.539**	**0.551**	0.303	0.404	**0.294**	**0.397**
	720	2.547	1.317	**1.402**	**0.949**	1.344	0.922	**1.010**	**0.772**	0.826	0.693	**0.819**	**0.654**
	Avg.	1.520	0.985	**0.814**	**0.712**	0.755	0.649	**0.527**	**0.525**	0.342	0.396	**0.337**	**0.384**
	Δ			-46.4%	-27.8%			-30.1%	-19.2%			-1.46%	-3.10%
Weather	96	0.203	0.287	**0.157**	**0.198**	0.145	0.209	**0.138**	**0.193**	0.150	0.196	**0.147**	**0.185**
	192	0.303	0.369	**0.216**	**0.256**	0.190	0.256	**0.184**	**0.245**	0.195	0.240	**0.190**	**0.225**
	336	0.351	0.373	**0.270**	**0.305**	0.252	0.306	**0.231**	**0.285**	0.246	0.280	**0.243**	**0.267**
	720	0.491	0.465	**0.448**	**0.435**	0.318	0.363	**0.298**	**0.336**	0.321	0.332	**0.319**	**0.321**
	Avg.	0.337	0.374	**0.273**	**0.299**	0.226	0.284	**0.213**	**0.265**	0.228	0.262	**0.225**	**0.250**
	Δ			-19.1%	-20.1%			-5.97%	-6.61%			-1.32%	-4.77%
ILI	24	4.689	1.466	**4.587**	**1.434**	3.595	1.265	**3.006**	**1.150**	1.900	0.868	**1.755**	**0.856**
	36	4.812	1.529	**3.273**	**1.243**	3.977	1.350	**3.416**	**1.234**	2.396	0.964	**2.056**	**0.926**
	48	4.952	1.572	**3.721**	**1.328**	3.783	1.297	**3.773**	**1.306**	1.938	0.917	**1.793**	**0.888**
	60	5.358	1.608	**4.405**	**1.434**	4.571	1.457	**4.527**	**1.450**	2.070	0.933	**2.014**	**0.911**
	Avg.	4.953	1.544	**3.997**	**1.360**	3.982	1.342	**3.681**	**1.285**	2.076	0.921	**1.905**	**0.895**
	Δ			-19.3%	-11.9%			-7.56%	-4.26%			-8.26%	-2.74%

Table 11: Part 2 of the full forecasting results without/with selective learning (SL). Better results are in **bold**. Avg. denotes the average from all prediction lengths, and Δ denotes the averaged improvements caused by selective learning.

Method		TimesNet				iTransformer				TimeMixer			
		w/o		+SL		w/o		+SL		w/o		+SL	
Metric		MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE
ETTh1	96	0.445	0.448	**0.398**	**0.411**	0.402	0.413	**0.371**	**0.389**	0.394	0.411	**0.367**	**0.387**
	192	0.476	0.472	**0.414**	**0.426**	0.445	0.440	**0.414**	**0.420**	0.440	0.442	**0.410**	**0.415**
	336	0.505	0.485	**0.440**	**0.446**	0.469	0.464	**0.431**	**0.432**	0.452	0.446	**0.434**	**0.429**
	720	0.571	0.537	**0.463**	**0.474**	0.514	0.510	**0.444**	**0.460**	0.485	0.480	**0.434**	**0.452**
	Avg.	0.499	0.486	**0.429**	**0.439**	0.458	0.457	**0.415**	**0.425**	0.443	0.445	**0.411**	**0.421**
	Δ			**-14.0%**	**-9.67%**			**-9.29%**	**-6.90%**			**-7.11%**	**-5.40%**
ETTh2	96	0.356	0.404	**0.292**	**0.357**	0.319	0.372	**0.300**	**0.352**	0.325	0.375	**0.299**	**0.354**
	192	0.427	0.452	**0.351**	**0.396**	0.394	0.419	**0.375**	**0.405**	0.412	0.436	**0.378**	**0.402**
	336	0.450	0.467	**0.389**	**0.423**	0.429	0.445	**0.408**	**0.433**	0.430	0.451	**0.405**	**0.428**
	720	0.505	0.500	**0.439**	**0.459**	0.460	0.474	**0.445**	**0.467**	0.457	0.472	**0.443**	**0.462**
	Avg.	0.435	0.456	**0.368**	**0.409**	0.401	0.428	**0.382**	**0.414**	0.406	0.434	**0.381**	**0.412**
	Δ			**-15.4%**	**-10.3%**			**-4.62%**	**-3.27%**			**-6.10%**	**-5.07%**
ETTm1	96	0.329	0.375	**0.298**	**0.342**	0.305	0.358	**0.295**	**0.342**	0.298	0.350	**0.287**	**0.337**
	192	0.377	0.402	**0.344**	**0.372**	0.346	0.380	**0.338**	**0.368**	0.329	0.370	**0.327**	**0.362**
	336	0.413	0.427	**0.382**	**0.397**	0.385	0.403	**0.371**	**0.386**	0.368	0.391	0.368	**0.382**
	720	0.464	0.453	**0.419**	**0.424**	0.446	0.441	**0.423**	**0.422**	0.431	0.426	**0.422**	**0.416**
	Avg.	0.396	0.414	**0.361**	**0.384**	0.371	0.396	**0.357**	**0.380**	0.357	0.384	**0.351**	**0.374**
	Δ			**-8.84%**	**-7.36%**			**-3.71%**	**-4.05%**			**-1.54%**	**-2.60%**
ETTm2	96	0.191	0.280	**0.169**	**0.257**	0.175	0.268	**0.166**	**0.253**	0.172	0.261	**0.164**	**0.250**
	192	0.246	0.314	**0.226**	**0.298**	0.244	0.315	**0.220**	**0.291**	0.233	0.305	**0.220**	**0.289**
	336	0.312	0.360	**0.271**	**0.324**	0.291	0.343	**0.270**	**0.325**	0.283	0.335	**0.267**	**0.321**
	720	0.408	0.416	**0.367**	**0.385**	0.383	0.400	**0.367**	**0.384**	0.370	0.391	**0.354**	**0.377**
	Avg.	0.289	0.343	**0.258**	**0.316**	0.273	0.332	**0.256**	**0.313**	0.265	0.323	**0.251**	**0.309**
	Δ			**-10.7%**	**-7.74%**			**-6.40%**	**-5.51%**			**-5.01%**	**-4.26%**
Eletricity	96	0.184	0.289	**0.177**	**0.281**	0.134	0.227	**0.132**	**0.223**	0.135	0.229	**0.133**	**0.228**
	192	0.192	0.295	**0.184**	**0.286**	0.157	0.249	**0.151**	**0.244**	0.152	0.247	**0.149**	**0.241**
	336	0.193	0.299	**0.190**	**0.294**	0.168	0.262	**0.158**	**0.250**	0.164	0.263	**0.160**	**0.255**
	720	0.222	0.320	**0.214**	**0.313**	0.197	0.290	**0.187**	**0.279**	0.201	0.297	**0.196**	**0.290**
	Avg.	0.198	0.301	**0.191**	**0.294**	0.164	0.257	**0.157**	**0.249**	0.163	0.259	**0.160**	**0.254**
	Δ			**-3.29%**	**-2.33%**			**-4.27%**	**-2.92%**			**-2.15%**	**-2.12%**
Exchange	96	0.109	0.239	**0.091**	**0.218**	0.088	0.211	**0.082**	**0.203**	0.080	0.199	**0.078**	**0.197**
	192	0.182	0.312	**0.165**	**0.299**	0.173	0.303	**0.162**	**0.293**	0.166	0.298	**0.157**	**0.287**
	336	0.333	0.427	**0.322**	**0.416**	0.323	0.419	**0.305**	**0.407**	0.310	0.407	**0.297**	**0.399**
	720	0.904	0.736	**0.875**	**0.720**	0.870	0.720	**0.821**	**0.693**	0.834	0.694	**0.808**	**0.693**
	Avg.	0.382	0.429	**0.363**	**0.413**	0.364	0.413	**0.343**	**0.399**	0.348	0.400	**0.335**	**0.394**
	Δ			**-4.97%**	**-3.73%**			**-5.78%**	**-3.45%**			**-3.60%**	**-1.38%**
Weather	96	0.167	0.222	**0.160**	**0.207**	0.161	0.210	**0.153**	**0.194**	0.152	0.207	0.152	**0.204**
	192	0.218	0.266	**0.212**	**0.255**	0.206	0.249	**0.197**	**0.237**	0.197	0.253	**0.194**	**0.243**
	336	0.265	0.300	**0.257**	**0.289**	0.249	0.280	**0.243**	**0.270**	0.249	0.293	**0.245**	**0.287**
	720	0.340	0.348	**0.326**	**0.334**	0.325	0.336	**0.321**	**0.328**	0.321	0.350	**0.312**	**0.338**
	Avg.	0.248	0.284	**0.239**	**0.271**	0.235	0.269	**0.229**	**0.257**	0.230	0.276	**0.226**	**0.268**
	Δ			**-3.54%**	**-4.58%**			**-2.87%**	**-4.28%**			**-1.74%**	**-2.81%**
ILI	24	2.480	1.009	**1.969**	**0.907**	1.511	0.813	**1.359**	**0.784**	1.931	0.879	**1.829**	**0.847**
	36	2.815	1.095	**2.405**	**0.979**	1.929	0.929	**1.696**	**0.847**	2.430	0.971	**2.113**	**0.902**
	48	2.436	1.008	**2.351**	**0.941**	2.054	0.931	**1.857**	**0.892**	2.135	0.931	**2.051**	**0.903**
	60	2.240	0.998	**1.892**	**0.898**	2.140	0.983	**1.926**	**0.905**	2.157	0.948	**2.109**	**0.926**
	Avg.	2.493	1.028	**2.154**	**0.931**	1.909	0.914	**1.710**	**0.857**	2.163	0.932	**2.026**	**0.895**
	Δ			**-13.6%**	**-9.36%**			**-10.4%**	**-6.24%**			**-6.37%**	**-4.06%**

23

Table 12: Full comparison results between selective learning (SL) and other training objectives with iTransformer as backbone. Avg. denotes the averaged results from all prediction lengths. The best results are in **bold**, and the second-best are underlined.

Loss		SL		PS		FreDF		TILDE-Q		MSE	
Metric		MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE
ETTh1	96	**0.371**	**0.389**	0.389	0.410	<u>0.388</u>	0.410	0.391	<u>0.408</u>	0.402	0.413
	192	**0.414**	**0.420**	<u>0.421</u>	0.429	0.434	0.438	0.423	<u>0.428</u>	0.445	0.440
	336	**0.431**	**0.432**	<u>0.443</u>	0.446	0.460	0.459	0.448	<u>0.444</u>	0.469	0.464
	720	**0.444**	**0.460**	<u>0.453</u>	<u>0.476</u>	0.516	0.513	0.467	0.477	0.514	0.510
	Avg.	**0.415**	**0.425**	<u>0.427</u>	0.440	0.450	0.455	0.432	<u>0.439</u>	0.458	0.457
ETTm2	96	**0.166**	**0.253**	<u>0.167</u>	<u>0.255</u>	0.169	0.258	0.173	0.259	0.175	0.268
	192	**0.220**	**0.291**	0.232	0.299	<u>0.227</u>	<u>0.298</u>	0.231	<u>0.298</u>	0.244	0.315
	336	**0.270**	**0.325**	0.287	0.336	<u>0.274</u>	<u>0.330</u>	0.278	0.331	0.291	0.343
	720	**0.367**	**0.384**	<u>0.370</u>	0.389	0.377	0.391	<u>0.370</u>	<u>0.388</u>	0.383	0.400
	Avg.	**0.256**	**0.313**	0.264	0.320	<u>0.262</u>	<u>0.319</u>	0.263	<u>0.319</u>	0.273	0.332
Exchange	96	**0.082**	**0.203**	0.087	0.211	0.088	0.208	<u>0.084</u>	<u>0.207</u>	0.088	0.211
	192	**0.162**	**0.293**	0.180	0.303	0.185	0.305	<u>0.171</u>	<u>0.301</u>	0.173	0.303
	336	**0.305**	**0.407**	0.335	0.420	0.346	0.426	0.335	0.422	<u>0.323</u>	<u>0.419</u>
	720	**0.821**	**0.693**	<u>0.861</u>	<u>0.700</u>	0.886	0.712	0.884	0.724	0.870	0.720
	Avg.	**0.343**	**0.399**	0.366	<u>0.409</u>	0.376	0.413	0.369	0.414	<u>0.364</u>	0.413
Weather	96	**0.153**	**0.194**	<u>0.155</u>	0.199	0.159	0.207	<u>0.155</u>	<u>0.198</u>	0.161	0.210
	192	**0.197**	**0.237**	0.200	0.241	0.204	0.249	<u>0.199</u>	<u>0.239</u>	0.206	0.249
	336	**0.243**	**0.270**	0.250	0.281	0.260	0.292	<u>0.249</u>	<u>0.279</u>	<u>0.249</u>	0.280
	720	**0.321**	**0.328**	0.327	0.337	0.334	0.347	<u>0.325</u>	<u>0.333</u>	<u>0.325</u>	0.336
	Avg.	**0.229**	**0.257**	0.233	0.265	0.239	0.274	<u>0.232</u>	<u>0.262</u>	0.235	0.269

Table 13: Full ablation results for selective learning with iTransformer as backbone.

| Design	Prediction Length	ETTh1		ETTm2		Electricity		Weather	
		MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE
Selective Learning	96	**0.371**	**0.389**	**0.166**	**0.253**	**0.132**	**0.223**	**0.153**	**0.194**
	192	**0.414**	**0.420**	**0.220**	**0.291**	**0.151**	**0.244**	**0.197**	**0.237**
	336	**0.431**	**0.432**	**0.270**	**0.325**	**0.158**	**0.250**	**0.243**	**0.270**
	720	**0.444**	**0.460**	**0.367**	**0.384**	**0.187**	**0.279**	**0.321**	**0.328**
	Avg.	**0.415**	**0.425**	**0.257**	**0.315**	**0.157**	**0.249**	**0.229**	**0.257**
w/o Uncertainty Mask	96	0.379	0.400	**0.166**	0.257	0.133	0.228	0.159	0.207
	192	0.423	0.429	0.230	0.300	0.153	0.247	0.201	0.246
	336	0.448	0.447	0.287	0.335	0.165	0.261	0.246	0.280
	720	0.495	0.497	0.377	0.394	0.195	0.287	0.321	0.331
	Avg.	0.436	0.443	0.265	0.322	0.162	0.256	0.232	0.266
w/o Anomaly Mask	96	0.388	0.403	0.177	0.261	0.135	0.228	0.158	0.206
	192	0.424	0.428	0.230	0.301	**0.151**	**0.244**	0.201	0.245
	336	0.444	0.442	0.279	0.333	0.160	0.256	0.249	0.281
	720	0.468	0.477	0.376	0.396	0.189	**0.279**	0.326	0.334
	Avg.	0.431	0.438	0.266	0.323	0.159	0.252	0.234	0.267
Random Mask	96	0.401	0.418	0.179	0.271	0.137	0.234	0.160	0.208
	192	0.452	0.452	0.240	0.312	0.158	0.256	0.203	0.249
	336	0.471	0.466	0.294	0.346	0.169	0.266	0.258	0.284
	720	0.505	0.505	0.384	0.399	0.196	0.286	0.325	0.334
	Avg.	0.457	0.460	0.274	0.332	0.165	0.261	0.237	0.269

(a) Case study on the ETTh1 datasets with iTransformer as backbone.

(b) Case study on the ETTm2 datasets with Informer as backbone.

(c) Case study on the Electricity datasets with Crossformer as backbone.

(d) Case study on the Exchange datasets with Informer as backbone.

(e) Case study on the Weather datasets with PatchTST as backbone.

Figure 6: Case study results across five datasets.

NeurIPS Paper Checklist

1. **Claims**

   Question: Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope?

   Answer: [Yes]

   Justification: The abstract and introduction clearly state the main contributions and scope of the paper.

   Guidelines:

   - The answer NA means that the abstract and introduction do not include the claims made in the paper.
   - The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A No or NA answer to this question will not be perceived well by the reviewers.
   - The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
   - It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. **Limitations**

   Question: Does the paper discuss the limitations of the work performed by the authors?

   Answer: [Yes]

   Justification: The paper discusses the limitations in the conclusion and appendix.

   Guidelines:

   - The answer NA means that the paper has no limitation while the answer No means that the paper has limitations, but those are not discussed in the paper.
   - The authors are encouraged to create a separate "Limitations" section in their paper.
   - The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
   - The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
   - The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
   - The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
   - If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
   - While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. **Theory assumptions and proofs**

   Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

   Answer: [Yes]

Justification: The paper provides the assumptions and proofs in the Appendix.

Guidelines:

- The answer NA means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. **Experimental result reproducibility**

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: The paper provides detailed descriptions of the datasets, experimental setups, implementation details, and hyperparameters, which are sufficient to reproduce the main results.

Guidelines:

- The answer NA means that the paper does not include experiments.
- If the paper includes experiments, a No answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general. releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
  (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
  (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
  (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
  (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. **Open access to data and code**

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: The paper provides the code in the supplemental material with sufficient instructions to reproduce the main experimental results. The code will be released publicly upon acceptance of the paper.

Guidelines:

- The answer NA means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (https://nips.cc/public/guides/CodeSubmissionPolicy) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so "No" is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (https://nips.cc/public/guides/CodeSubmissionPolicy) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. **Experimental setting/details**

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer, etc.) necessary to understand the results?

Answer: [Yes]

Justification: The paper thoroughly describes the training and test details.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. **Experiment statistical significance**

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: The paper provides clear explanations of the statistical methods used to compute metrics and ensures the robustness of the reported findings.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The authors should answer "Yes" if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.

- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g. negative error rates).
- If error bars are reported in tables or plots, The authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. **Experiments compute resources**

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: The paper provide sufficient information on the computer resources.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. **Code of ethics**

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics https://neurips.cc/public/EthicsGuidelines?

Answer: [Yes]

Justification: The research aligns with the NeurIPS Code of Ethics.

Guidelines:

- The answer NA means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer No, they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. **Broader impacts**

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [NA]

Justification: The paper focuses on time series forecasting problem, so there is no potential societal impact.

Guidelines:

- The answer NA means that there is no societal impact of the work performed.

- If the authors answer NA or No, they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. **Safeguards**

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pretrained language models, image generators, or scraped datasets)?

Answer: [NA]

Justification: The paper poses no such risks.

Guidelines:

- The answer NA means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. **Licenses for existing assets**

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: The paper credits the owners of any used assets.

Guidelines:

- The answer NA means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.

- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, `paperswithcode.com/datasets` has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. **New assets**

    Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

    Answer: [Yes]

    Justification: The paper introduces new assets that are well documented alongside the assets.

    Guidelines:

    - The answer NA means that the paper does not release new assets.
    - Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
    - The paper should discuss whether and how consent was obtained from people whose asset is used.
    - At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. **Crowdsourcing and research with human subjects**

    Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

    Answer: [NA]

    Justification: This paper does not involve crowdsourcing nor research with human subjects.

    Guidelines:

    - The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
    - Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
    - According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. **Institutional review board (IRB) approvals or equivalent for research with human subjects**

    Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

    Answer: [NA]

    Justification: The research does not involve human subjects, so IRB approvals or equivalent review are not required.

    Guidelines:

    - The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
    - Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.

- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. **Declaration of LLM usage**

   Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does not impact the core methodology, scientific rigorousness, or originality of the research, declaration is not required.

   Answer: [NA]

   Justification: The core method in this research does not involve LLMs as any important, original, or non-standard components.

   Guidelines:

   - The answer NA means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
   - Please refer to our LLM policy (`https://neurips.cc/Conferences/2025/LLM`) for what should or should not be described.